

CSC 433/533 Computer Graphics

Lecture 06 Anti-Aliasing and Signal Processing

Recall:
Images are Functions

Domains and Ranges

- All functions have two components, the **domain** and **range**. For the case of images, $I: R \rightarrow V$
- The domain is:
 - R , is some rectangular area ($R \subseteq \mathbb{R}^2$)
- The range is:
 - A set of possible values.
 - ...in the space of color values we're encoding



Operations on Images

Slides inspired from Fredo Durand

- Point (Range) Operations:



- Affect only the range of the image (e.g. brightness)
- Each pixel is processed separately, only depending on the color

Operations on Images

Slides inspired from Fredo Durand

- Domain Operations:



- Only move the pixels around

Operations on Images

Slides inspired from Fredo Durand

- Neighborhood operations:



- Combine domain and range
- Each pixel evaluated by working with other pixels nearby

Concept for the Day: Pixels are Samples of Image Functions

Image Samples

- Each pixel is a sample of what?
 - One interpretation: a pixel represents the intensity of light at a single (infinitely small point in space)
- The sample is displayed in such a way as to spread the point out across some spatial area (drawing a square of color)

Continuous vs. Discrete

- Key Idea: An image represents data in either (both?) of
 - Continuous domain: where light intensity is defined at every (infinitesimally small) point in some projection
 - Discrete domain, where intensity is defined only at a discretely sampled set of points.

Converting Between Image Domains

- When an image is acquired, an image is **sampling** from some continuous domain to a discrete domain.
- **Reconstruction** converts digital back to continuous.
- The reconstructed image can then be **resampled** and **quantized** back to the discrete domain.

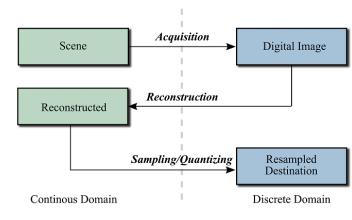


Figure 7.7. Resampling.

```

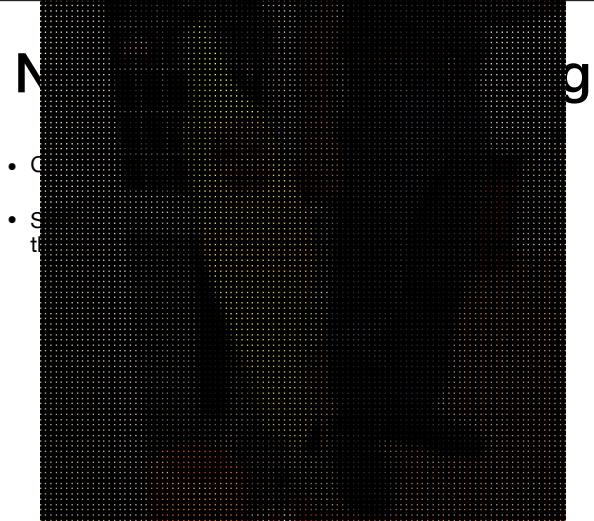
//scale factor
let k = 4;

//create an output greyscale image that is both
//k times as wide and k times as tall
Uint8Array output = new Uint8Array((k*W)*(k*H));

//copy the pixels over
for (let row = 0, row < H; row++) {
    for (let col = 0; col < W; col++) {
        let index = row*W + col;
        let index2 = (k*row)*W + (k*col);
        output[index2] = input[index];
    }
}

```

Naive Image Rescaling Code



What's the Problem?

- The output image has gaps!
- Why: we skip a many of the pixels in the output.
- Why don't we fix this by changing the code to at least put some color at each pixel of the output?

Naive Image Rescaling Code

```

//scale factor
let k = 4;

//create an output greyscale image that is both
//k times as wide and k times as tall
Uint8Array output = new Uint8Array((k*W)*(k*H));

//copy the pixels over
for (let row = 0, row < H; row++) {
    for (let col = 0; col < W; col++) {
        let index = row*W + col;
        let index2 = (k*row)*W + (k*col);
        output[index2] = input[index];
    }
}

```

```

//scale factor
let k = 4;

//create an output greyscale image that is both
//k times as wide and k times as tall
Uint8Array output = new Uint8Array((k*W)*(k*H));

//Loop over each output pixel instead.
for (let row = 0, row < k*H; row++) {
    for (let col = 0; col < k*W; col++) {
        let index = (row/k)*W + (col/k);
        let index2 = row*k*W + col;
        output[index2] = input[index];
    }
}

```

"Inverse" Image Rescaling Code



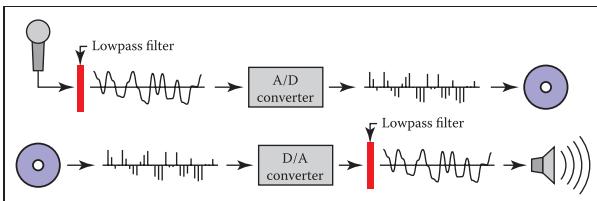
What's the Problem?

- The output image is too “blocky”
- Why: because our image reconstruction rounds the index to the nearest integer pixel coordinates
- Rounding to the “nearest” is why this type of interpolation is called **nearest neighbor interpolation**

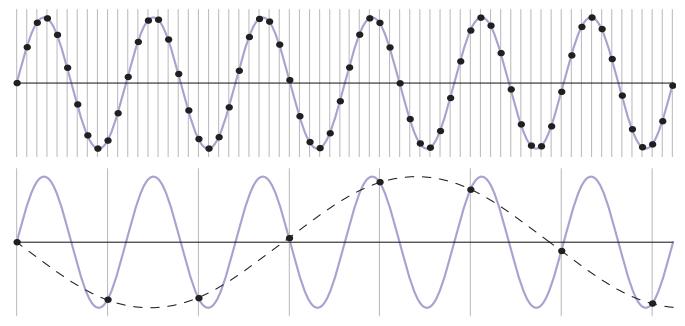
Sampling Artifacts / Aliasing

Motivation: Digital Audio

- Acquisition of images takes a continuous object and converts this signal to something digital
- Two types of artifacts:
 - **Undersampling** artifacts: on acquisition side
 - **Reconstruction** artifacts: when the samples are interpreted



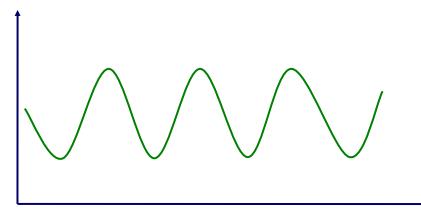
Undersampling Artifacts



S-N Theorem Illustrated

How many samples are enough to avoid aliasing?

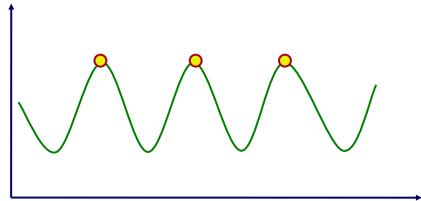
- How many samples are required to represent a given signal without loss of information?
- What signals can be reconstructed without loss for a given sampling rate?



S-N Theorem Illustrated

How many samples are enough to avoid aliasing?

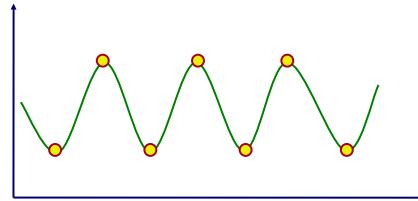
- How many samples are required to represent a given signal without loss of information?
- What signals can be reconstructed without loss for a given sampling rate?



S-N Theorem Illustrated

How many samples are enough to avoid aliasing?

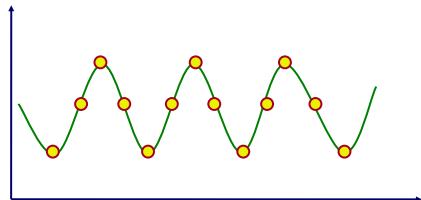
- How many samples are required to represent a given signal without loss of information?
- What signals can be reconstructed without loss for a given sampling rate?



S-N Theorem Illustrated

How many samples are enough to avoid aliasing?

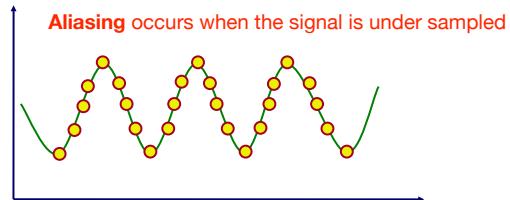
- How many samples are required to represent a given signal without loss of information?
- What signals can be reconstructed without loss for a given sampling rate?



S-N Theorem Illustrated

How many samples are enough to avoid aliasing?

- How many samples are required to represent a given signal without loss of information?
- What signals can be reconstructed without loss for a given sampling rate?



Temporal Aliasing: The Wagon Wheel Effect

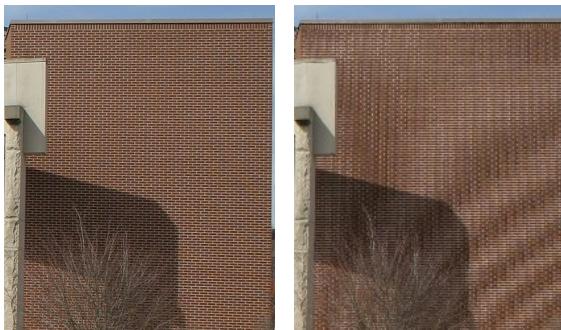


Aliasing in images

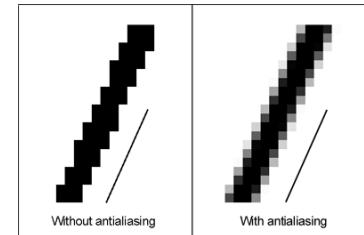
Two outcomes of under-sampling

- 1) Moire Pattern
- 2) Rasterization

Moire Patterns



Aliasing for edges



Each pixel is effected by nearby pixels
For example, even though the input image image is black/white,
We allow grey values for output pixels.

Convolution

Each pixel is effected by nearby pixels
For example, even though the image is black/white,
We allow grey values

An Example: Mean Filtering

- Mean filters sum all of the pixels in a local neighborhood N_i and divide by the total number, computing the average pixel.
- Said another way, we replace each pixel as a linear combination of its neighbors (with equal weights!)
- $f(N_i) = 1 / |N_i| \sum C_j$, for pixel j in N_i
- Where the N_i is a square, we call these **box** filters

Box Filtering



$$\begin{matrix} 0 & 0 & 0 \\ 0 & 1 & 0 \\ 0 & 0 & 0 \end{matrix}$$



$$1/9 * \begin{matrix} 1 & 1 & 1 \\ 1 & 1 & 1 \\ 1 & 1 & 1 \end{matrix}$$

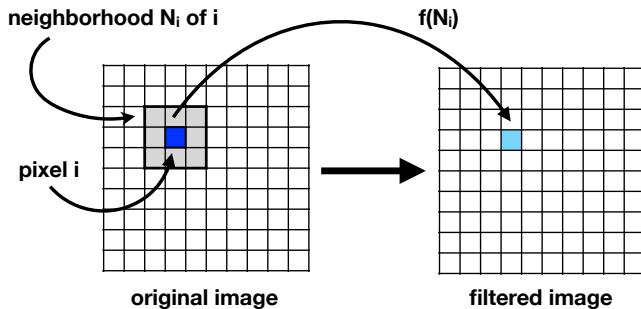
Box Filtering



$$1/25 * \begin{matrix} 1 & 1 & 1 & 1 & 1 \\ 1 & 1 & 1 & 1 & 1 \\ 1 & 1 & 1 & 1 & 1 \\ 1 & 1 & 1 & 1 & 1 \\ 1 & 1 & 1 & 1 & 1 \end{matrix}$$

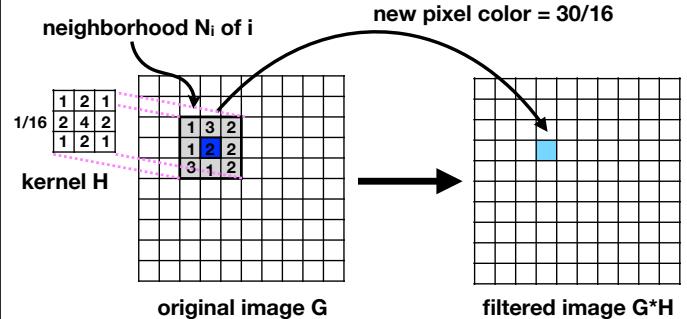


Neighborhood Filtering (Schematic)



Convolution

- This process of adding up pixels multiplied by various weights is called **convolution**



Kernels

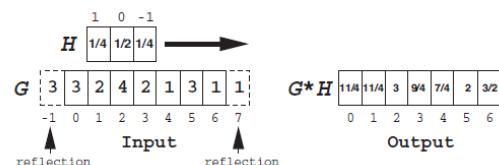
- Convolution employs a rectangular grid of coefficients, known as a **kernel**
- Kernels are like a neighborhood mask, they specify which elements of the image are in the neighborhood and their relative weights.
- A kernel is a set of weights that is applied to corresponding input samples that are summed to produce the output sample.

One-dimensional Convolution

- Can be expressed by the following equation, which takes a filter H and convolves it with G :

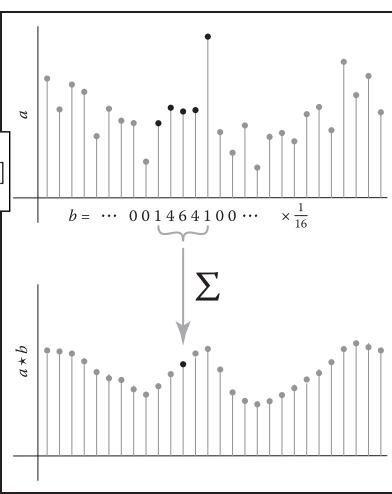
$$\hat{G}[i] = (G * H)[i] = \sum_{j=i-n}^{i+n} G[j]H[i-j], \quad i \in [0, N-1]$$

- Equivalent to sliding a window



Convolution is a Moving, Weighted Average

$$(a \star b)[i] = \sum_{j=i-r}^{i+r} a[j]b[i-j]$$



- Mathematically, this is equivalent to integrating the product of a and b with a shift in the domain

- Compare a to $a*b$ on the right

2-Dimensional Version

- Given an image a and a kernel b with $(2r+1)^2$ values, the convolution of a with b is given below as $a*b$:

$$(a \star b)[i, j] = \sum_{i'=i-r}^{i+r} \sum_{j'=j-r}^{j+r} a[i', j']b[i - i', j - j']$$

- The $(i-i')$ and $(j-j')$ terms can be understood as reflections of the kernel about the central vertical and horizontal axes.
- The kernel weights are multiplied by the corresponding image samples and then summed together.

A Note on Indexing

- Convolution **reflects** the filter to preserve orientation.
 - Correlation does **not** have this reflection.
 - But we often use them interchangeably since most kernels are symmetric!!

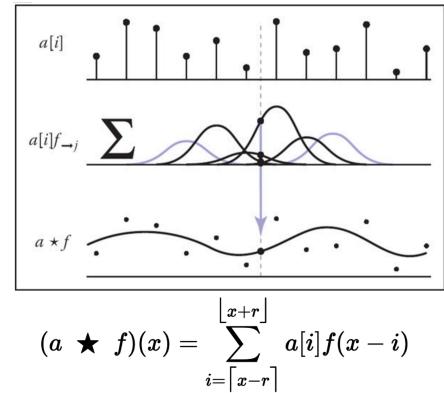
**Convolution reflects
and shifts the kernel**

Given kernel H =

1	2	3
4	5	6
7	8	9

Convolution Can Also Convert from Discrete to Continuous

- Discrete signal a
 - Continuous filter f
 - Output $a \cdot f$ defined on positions x as opposed to discrete pixels i



Filtering helps to reconstruct the signal better when rescaling



Inverse Rescaling

Reconstructed w/ Discrete-to-Continuous

```
//scale factor
let k = 4;

//create an output greyscale image that is both
//k times as wide and k times as tall
Uint8Array output = new Uint8Array((k*W)*(k*H));

//Loop over each output pixel instead.
for (let row = 0, row < k*H; row++) {
    for (let col = 0; col < k*W; col++) {
        let x = col/k;
        let y = row/k;
        let index = row*k*W + col;
        output[index] = reconstruct(input,x,y);
    }
}
```

Types of Filters: Smoothing

Smoothing Spatial Filters

- Any weighted filter with positive values will smooth in some way, examples:

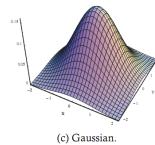
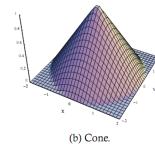
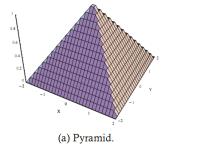
$$\frac{1}{9} \times \begin{array}{|c|c|c|} \hline 1 & 1 & 1 \\ \hline 1 & 1 & 1 \\ \hline 1 & 1 & 1 \\ \hline \end{array} \quad \frac{1}{16} \times \begin{array}{|c|c|c|} \hline 1 & 2 & 1 \\ \hline 2 & 4 & 2 \\ \hline 1 & 2 & 1 \\ \hline \end{array}$$

- Normally, we use integers in the filter, and then divide by the sum (computationally more efficient)
- These are also called **blurring** or **low-pass** filters

Smoothing Kernels

$$f(x, y) = -\alpha \cdot \max(|x|, |y|)$$

$$G(x, y) = \frac{1}{2\pi\sigma^2} e^{-\frac{(x^2+y^2)}{2\sigma^2}}$$



$$f(x, y) = -\alpha \cdot \sqrt{x^2 + y^2}$$

$$\begin{array}{|c|c|c|c|c|} \hline 1 & 2 & 3 & 2 & 1 \\ \hline 2 & 4 & 6 & 4 & 2 \\ \hline 3 & 6 & 9 & 6 & 3 \\ \hline 2 & 4 & 6 & 4 & 2 \\ \hline 1 & 2 & 3 & 2 & 1 \\ \hline \end{array}$$

(a) Pyramid.

$$\begin{array}{|c|c|c|c|c|} \hline 0 & 0 & 1 & 0 & 0 \\ \hline 0 & 2 & 2 & 2 & 0 \\ \hline 1 & 2 & 5 & 2 & 1 \\ \hline 0 & 2 & 2 & 2 & 0 \\ \hline 1 & 0 & 1 & 0 & 0 \\ \hline \end{array}$$

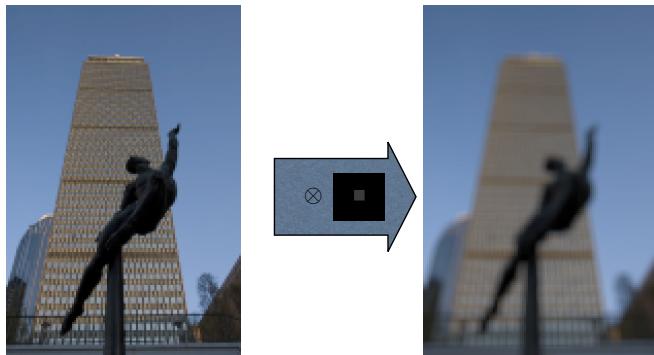
(b) Cone.

$$\begin{array}{|c|c|c|c|c|} \hline 1 & 4 & 7 & 4 & 1 \\ \hline 4 & 16 & 28 & 16 & 4 \\ \hline 7 & 28 & 49 & 28 & 7 \\ \hline 4 & 16 & 28 & 16 & 4 \\ \hline 1 & 4 & 7 & 4 & 1 \\ \hline \end{array}$$

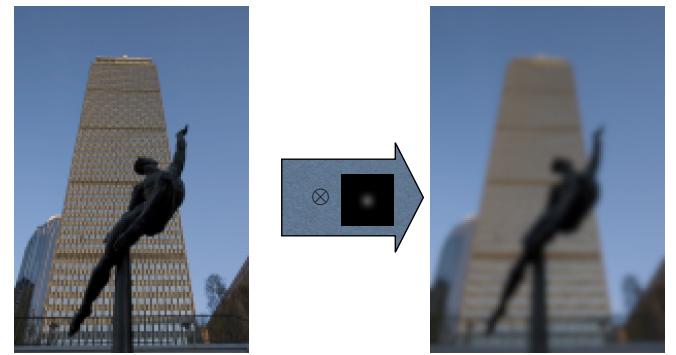
(c) Gaussian.

Table 6.1. Discretized kernels.

Box Filter



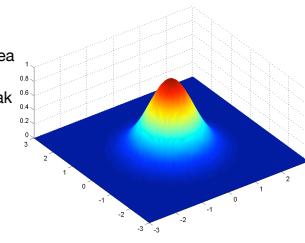
Gaussian Filter



Gaussians

$$G(x, y) = \frac{1}{2\pi\sigma^2} e^{-\frac{(x^2+y^2)}{2\sigma^2}}$$

- Gaussian kernel is parameterized on the standard deviation σ
- Large σ 's reduce the center peak and spread the information across a larger area
- Smaller σ 's create a thinner and taller peak
- Gaussians are smooth everywhere.
- Gaussians have infinite support
- >0 everywhere
- But often truncate to 2σ or 3σ



http://en.wikipedia.org/wiki/Gaussian_function

Smoothing Comparison

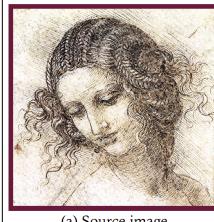
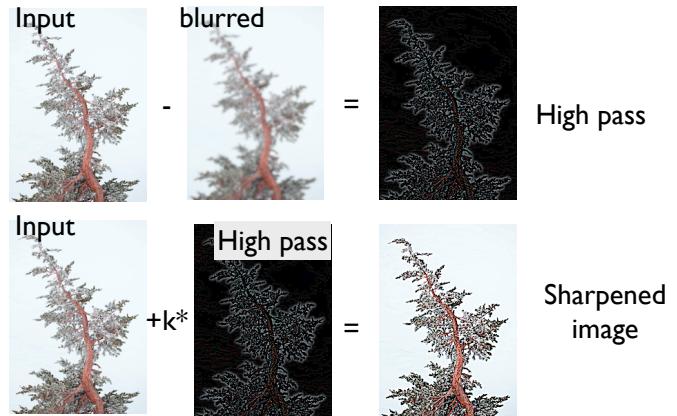


Figure 6.10. Smoothing examples.

Types of Filters: Sharpening

Sharpening (Idea)



Sharpening is a Convolution

- This procedure can then be expressed as a single kernel
- Assume that $I = I * d$ and $I_{\text{low}} = I * f_{g,\sigma}$.
 - d is the discrete identity function (kernel with 1 in center, 0 elsewhere)
 - $f_{g,\sigma}$ is a smoothing filter (e.g. Gaussian of width σ).
- This leads to:

$$\begin{aligned} I_{\text{sharp}} &= (1 + \alpha)I - \alpha(I * f_{g,\sigma}) \\ &= I * ((1 + \alpha)d - \alpha f_{g,\sigma}) \\ &= I * f_{\text{sharp}}(\sigma, \alpha), \end{aligned}$$

Sharpening is a Convolution

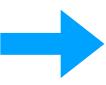
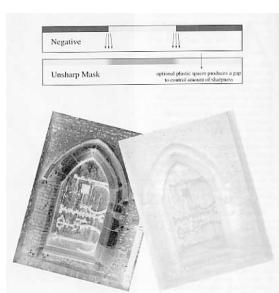
$$\begin{aligned} I_{\text{sharp}} &= (1 + \alpha)I - \alpha(I * f_{g,\sigma}) \\ &= I * ((1 + \alpha)d - \alpha f_{g,\sigma}) \\ &= I * f_{\text{sharp}}(\sigma, \alpha), \end{aligned}$$

Note: could also define d as $\begin{bmatrix} 0 & 0 & 0 \\ 0 & 1 & 0 \\ 0 & 0 & 0 \end{bmatrix}$

$$\begin{aligned} d &= \frac{1}{9} \times \begin{bmatrix} 0 & 0 & 0 \\ 0 & 9 & 0 \\ 0 & 0 & 0 \end{bmatrix}, \\ f_{g,\sigma} &= \frac{1}{9} \times \begin{bmatrix} 1 & 1 & 1 \\ 1 & 1 & 1 \\ 1 & 1 & 1 \end{bmatrix}, \\ ((1 + \alpha)d - \alpha f_{g,\sigma}) &= \frac{1}{9} \times \begin{bmatrix} -\alpha & -\alpha & -\alpha \\ -\alpha & (9 + 8\alpha) & -\alpha \\ -\alpha & -\alpha & -\alpha \end{bmatrix} \end{aligned}$$

Unsharp Masks

- Sharpening is often called “unsharp mask” because photographers used to sandwich a negative with a blurry positive film in order to sharpen



<http://www.tech-diy.com/UnsharpMasks.htm>

Edge Enhancement

- The parameter α controls how much of the source image is passed through to the sharpened image.



Figure 6.20. Image sharpening.

Defining Edges

- Sharpening uses negative weights to enhance regions where the image is changing rapidly
- These rapid transitions between light and dark regions are called **edges**
- Smoothing reduces the strength of edges, sharpening strengthens them.
- Also called **high-pass filters**
- Idea: smoothing filters are weighted averages, or integrals. Sharpening filters are weighted differences, or derivatives!

Edges

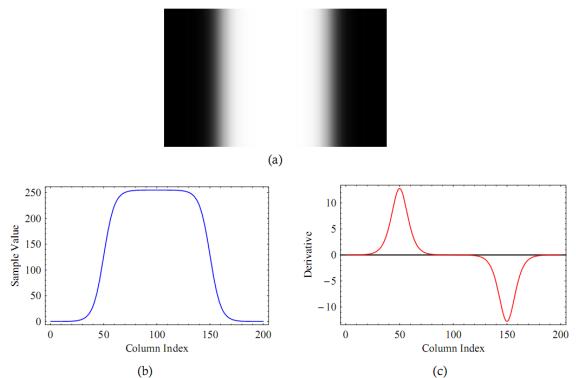


Figure 6.11. (a) A grayscale image with two edges, (b) row profile, and (c) first derivative.

(Review?) Derivatives via Finite Differences

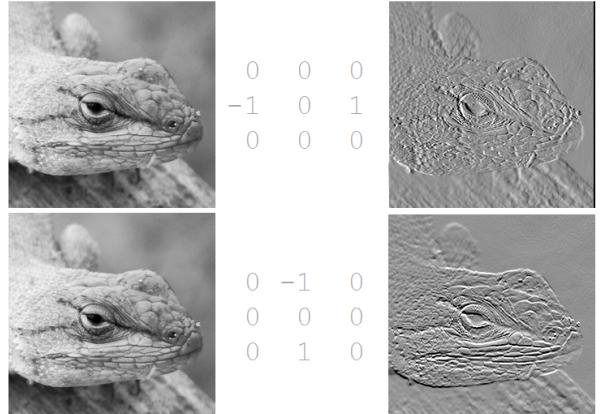
- We can approximate the derivative with a kernel w:

$$\frac{\partial f}{\partial x} \approx \frac{1}{2h} (f(x+1, y) - f(x-1, y))$$

$$\frac{\partial f}{\partial x} \approx w_{dx} \circ f \quad w_{dx} = \begin{bmatrix} -\frac{1}{2} & 0 & \frac{1}{2} \end{bmatrix}$$

$$\frac{\partial f}{\partial y} \approx w_{dy} \circ f \quad w_{dy} = \begin{bmatrix} -\frac{1}{2} \\ 0 \\ \frac{1}{2} \end{bmatrix}$$

Taking Derivatives with Convolution



Gradients with Finite Differences

- These partial derivatives approximate the image gradient, ∇I .
- Gradients are the unique direction where the image is changing the most rapidly, like a slope in high dimensions
- We can separate them into components kernels G_x, G_y . $\nabla I = (G_x, G_y)$

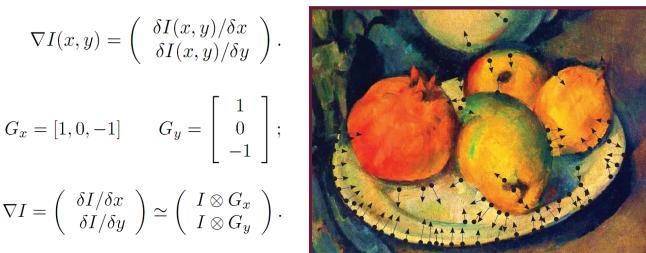


Figure 6.13. Image gradient (partial).

128	187	210	238	251
76	121	193	225	219
66	91	110	165	205
47	81	83	119	157
41	59	63	75	125

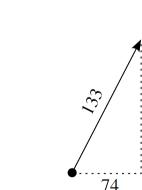
(a) Source Image.

117	104	26	
44	74	95	
36	38	74	

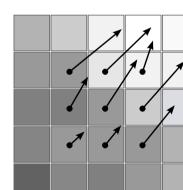
(b) $\delta I / \delta x$.

-96	-100	-73	
-40	-110	-106	
-32	-47	-90	

(c) $\delta I / \delta y$.



(d) Center sample gradient.



(e) Gradient.

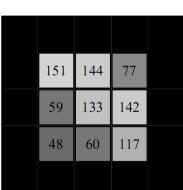
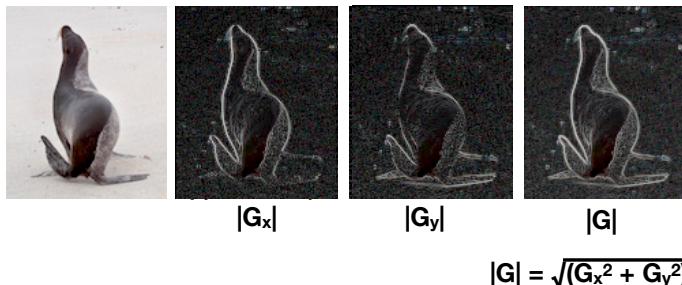


Figure 6.14. Numeric example of an image gradient.

Gradients G_x, G_y



Second Derivatives (Sharpening, almost)

- Partial derivatives in x and y lead to two kernels:

$$\frac{\partial^2 f}{\partial x^2} = f(x+1, y) + f(x-1, y) - 2f(x, y)$$

and, similarly, in the y-direction we have

$$\frac{\partial^2 f}{\partial y^2} = f(x, y+1) + f(x, y-1) - 2f(x, y)$$

0	1	0
1	-4	1
0	1	0

1	1	1
1	-8	1
1	1	1

Compare with
Sharpening filter:
unbalanced counts!

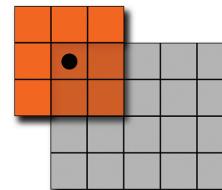
$$\begin{bmatrix} -\alpha & -\alpha & -\alpha \\ -\alpha & (9+8\alpha) & -\alpha \\ -\alpha & -\alpha & -\alpha \end{bmatrix}$$

1	1	1
1	-9	1
1	1	1

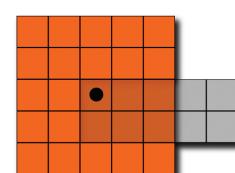
Boundaries

Handling Image Boundaries

- What should be done if the kernel falls off the boundary of the source image as shown in the illustrations below?



(a) Kernel at $I(0, 0)$.



(b) Kernel larger than the source.

Figure 6.4. Illustration of the edge handling problem.

Handling Image Boundaries

- When pixels are near the edge of the image, neighborhoods become tricky to define
- Choices:
 - Shrink the output image (ignore pixels near the boundary)
 - Expanding the input image (padding to create values near the boundary which are “meaningful”)
 - Shrink the kernel (skip values that are outside the boundary, and reweigh accordingly)

Boundary Padding

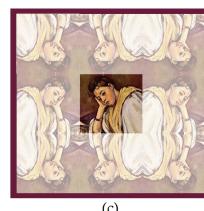
- When one pads, they pretend the image is large and either produce a constant (e.g. zero), or use circular / reflected indexing to tile the image:



(a)



(b)



(c)

Figure 6.5. (a) Zero padding, (b) circular indexing, and (c) reflected indexing.

Lec06 Required Reading

- FOCG, Ch. 21 (particularly 21.1-21.3)
- See also optional readings (especially grads!)

Reminder: Assignment 02

Assigned: Wednesday, Sept. 11
Written Due: Monday, Sept. 23, 4:59:59 pm
Program Due: Wednesday, Sept. 25, 4:59:59 pm