

Tone Reproduction

Today's Agenda

- Reminders:
 - A02 questions?
- Goals for today:
 - Wrap up filtering
 - Introduction point processing and tone reproduction

First Example: Linear Rescaling

- **Rescaling** is a point processing technique that alters the **contrast** and/or **brightness** of an image.
- In photography, **exposure** is a measure of how much light is projected onto the imaging sensor.
 - **Overexposure**: more light than what the sensor can measure.
 - **Underexposure**: sensor is unable to detect the light.
- Images which are underexposed or overexposed can frequently be improved by brightening or darkening them.
- The contrast of an image can be altered to bring out the internal structure of the image.

Rescaling Math

- Given a sample C_{in} of the source image, rescaling computes the output sample, C_{out} , using the scaling function

$$C_{out} = \alpha C_{in} + \beta$$

- α is a real-valued scaling factor known as **gain**
- β is a real-valued scaling factor known as **bias**

Effects of Rescaling

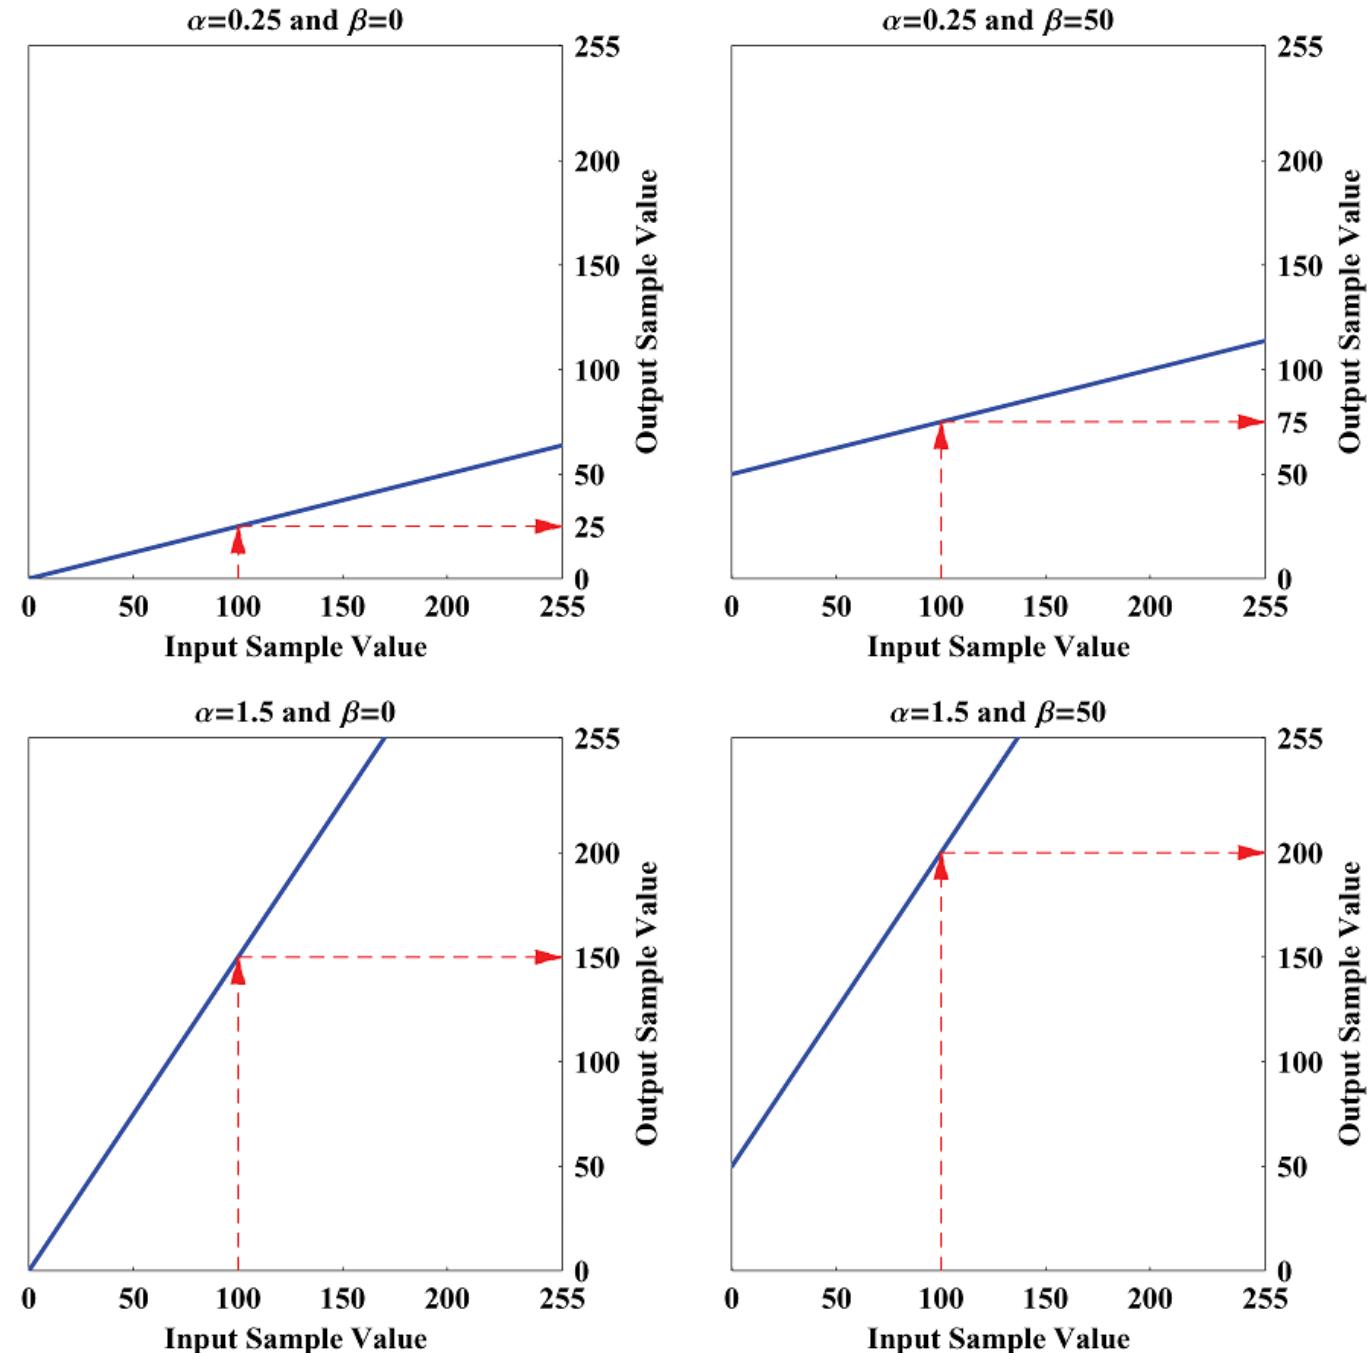


Figure 5.2. Graph of the linear scaling function with various gain and bias settings.

Why Use Both α , β ?

- Consider two rescaled source samples of S rescaled to S' .
- Calculate the **contrast** (the absolute difference) between the source and destination, called ΔS and $\Delta S'$.
- Now consider the relative change in contrast between the source and destination.

$$\begin{aligned}S'_1 &= \alpha S_1 + \beta, \\S'_2 &= \alpha S_2 + \beta.\end{aligned}$$

$$\begin{aligned}\Delta S' &= |S'_1 - S'_2|, \\ \Delta S &= |S_1 - S_2|.\end{aligned}$$

$$\frac{\Delta S'}{\Delta S} = \frac{|S'_1 - S'_2|}{|S_1 - S_2|}.$$

Why Use Both α , β ?

- The relative change in contrast can be simplified as

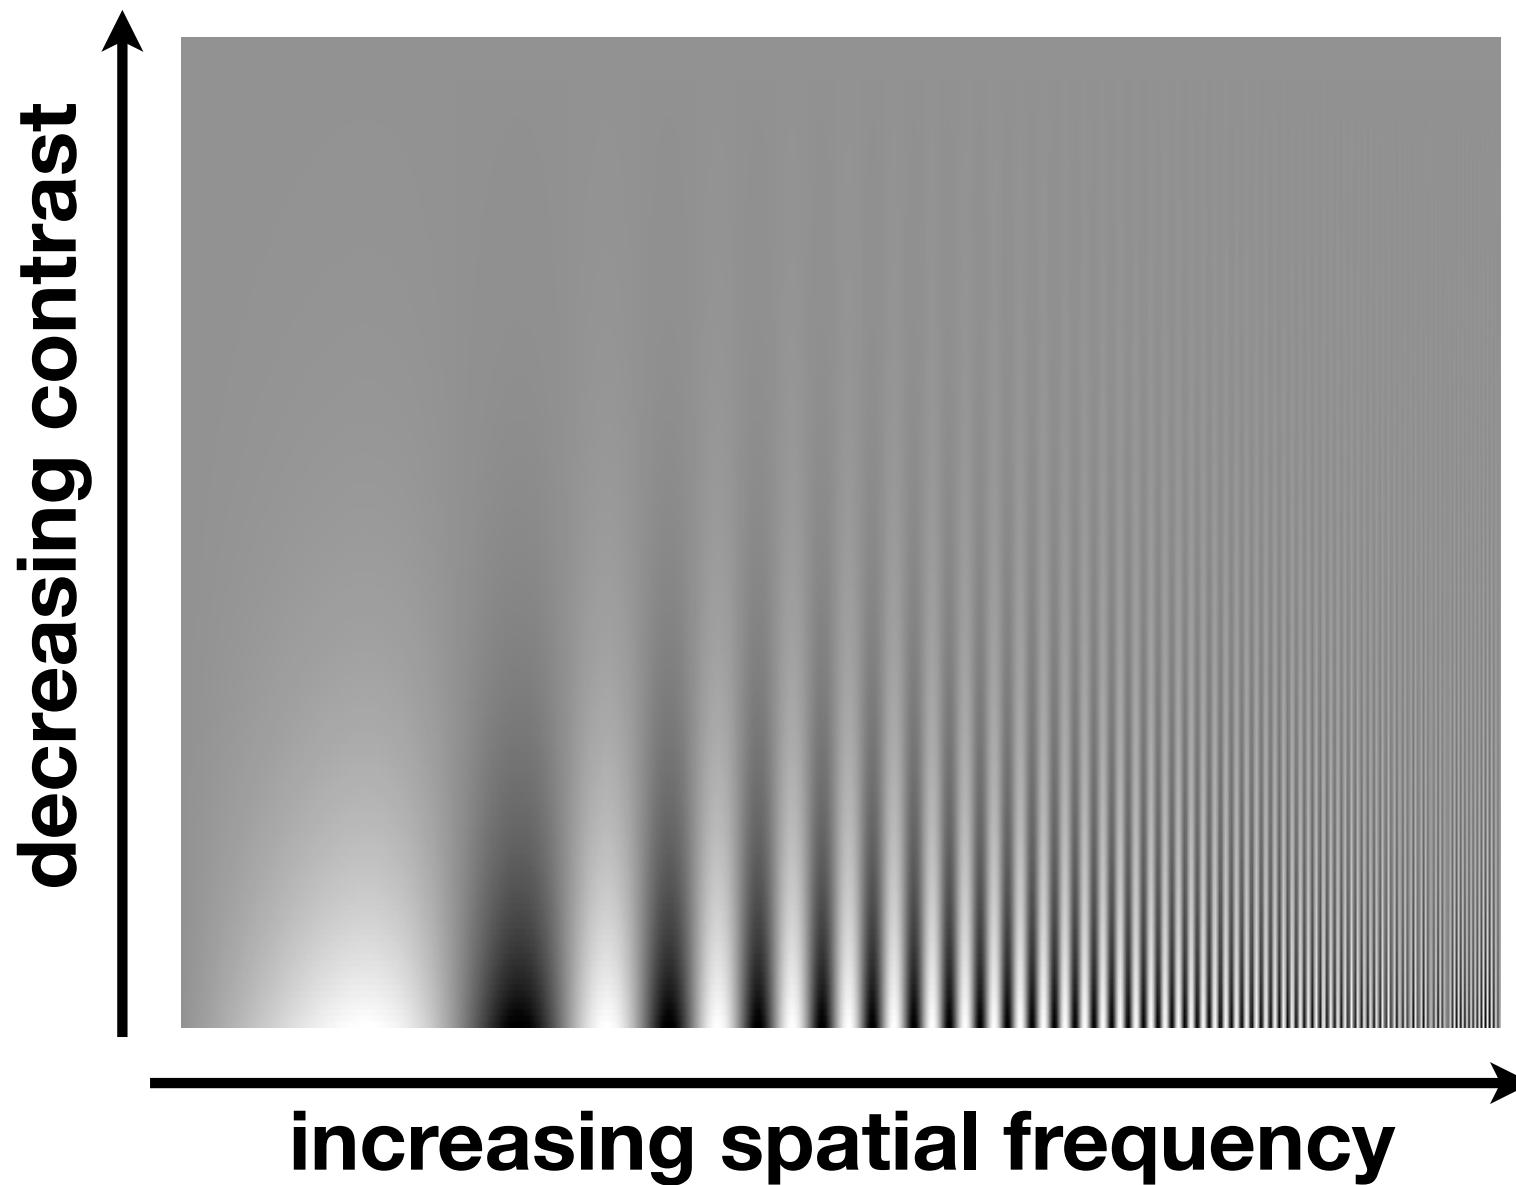
$$\begin{aligned}\frac{\Delta S'}{\Delta S} &= \frac{|(\alpha S_1 + \beta) - (\alpha S_2 + \beta)|}{|S_1 - S_2|} \\ &= \frac{|\alpha| \cdot |S_1 - S_2|}{|S_1 - S_2|} \\ &= |\alpha|.\end{aligned}$$

- Thus, gain (α) controls the change in contrast.
- Whereas bias (β) *does not* affect the contrast
- Bias, however, controls the final **brightness** of the rescaled image. Negative bias darkens and positive bias brightens the image

Sidebar: Relating Contrast Sensitivities to Signal Processing

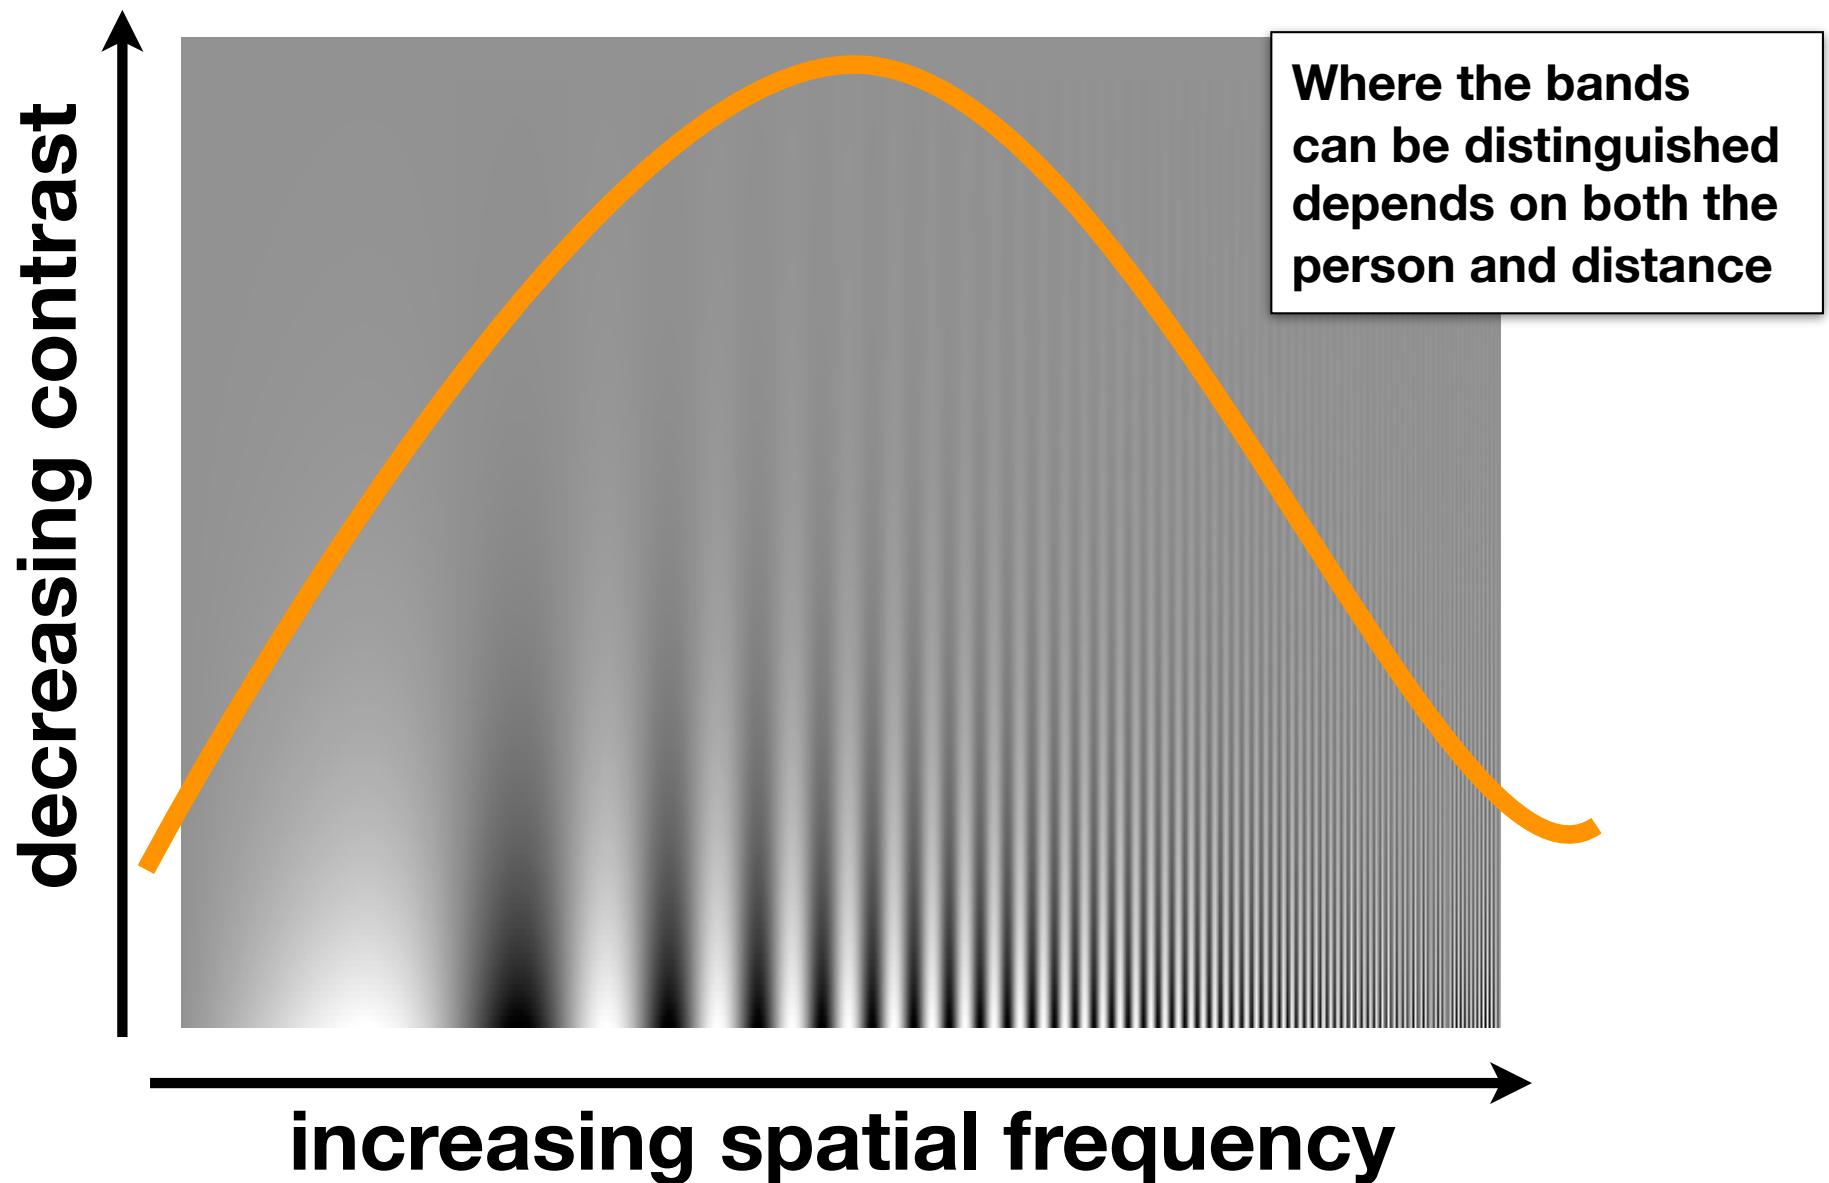
Contrast Sensitivity Function

Campbell-Robson Chart



Contrast Sensitivity Function

Campbell-Robson Chart



Contrast Sensitivities Vary by Channel

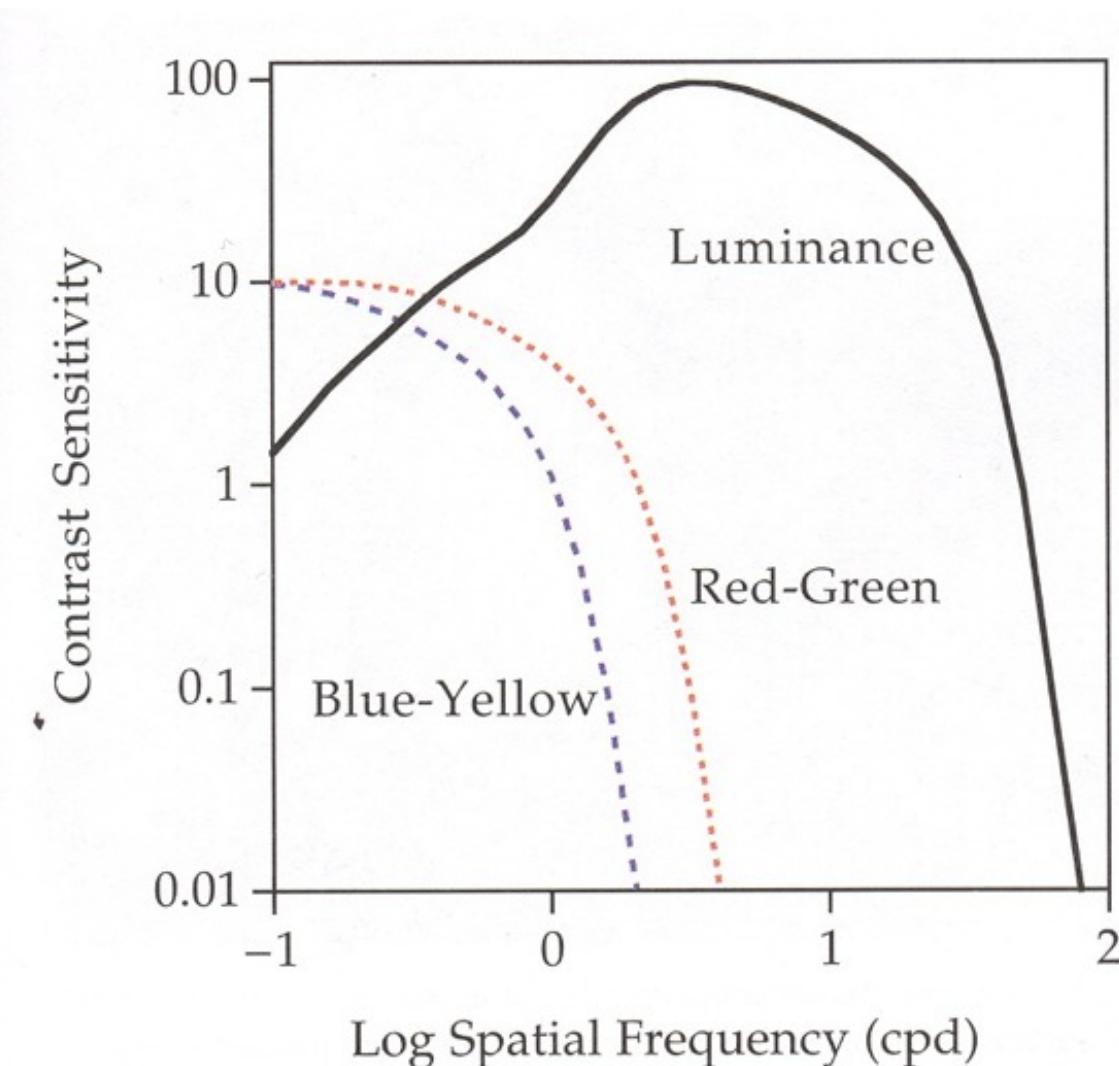


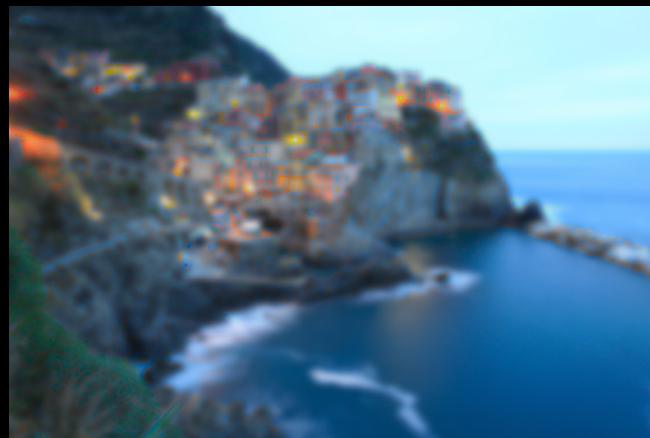
Figure 1-18. Spatial contrast sensitivity functions for luminance and chromatic contrast.

Photoshop demo

- Image > Mode > Lab color
- Go to channel panel, select Lightness
- Filter > Blur > Gaussian Blur , e.g. 4 pixel radius
 - very noticeable
- Undo, then select a & b channels
- Filter > Blur > Gaussian Blur , same radius
 - hardly visible effect



Original



Blur Lightness



Blur a & b

Important: Clamping

- Rescaling may produce samples that lie outside of the output images (e.g. below 0 or above 255 in 8-bit images)
- **Clamping** the output values ensures that the output samples are truncated to the 8-bit dynamic range limit
- Note that clamping does ‘lose’ information, since it truncates.

$$clamp(x, min, max) = \begin{cases} \min & \text{if } \lfloor x \rfloor \leq \min, \\ \max & \text{if } \lfloor x \rfloor \geq \max, \\ \lfloor x \rfloor & \text{otherwise.} \end{cases}$$

Rescaling Examples



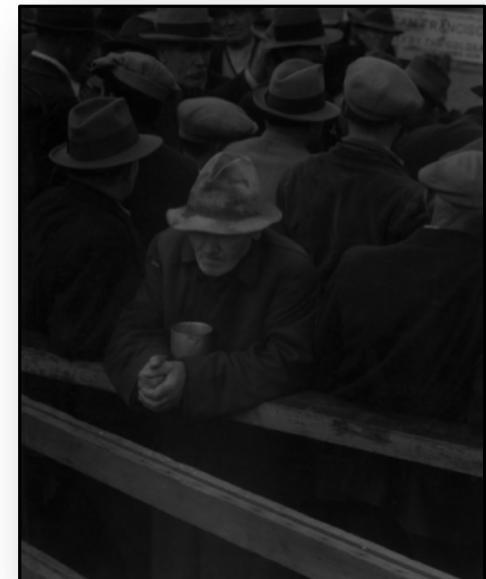
gain = 1, bias = 55



gain = 1, bias = -55



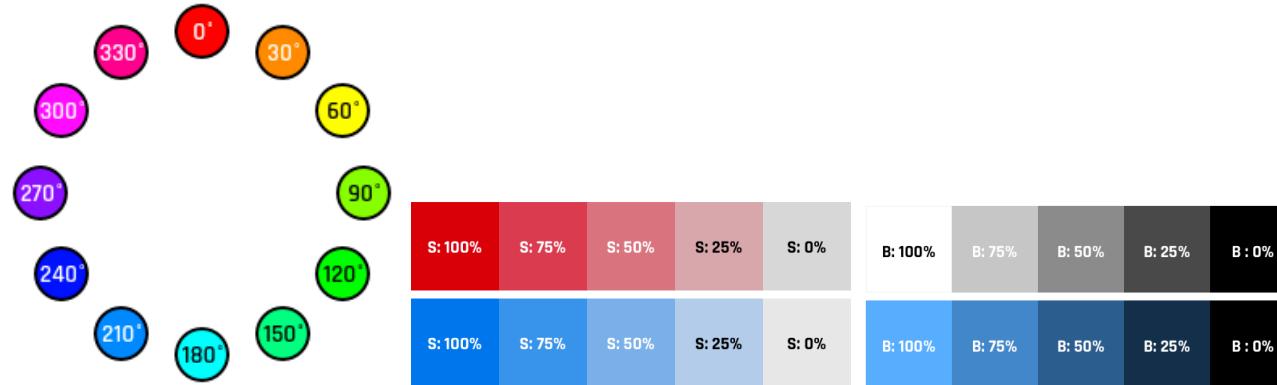
gain = 2, bias=0



gain = .5, bias=0

Rescaling Color Images

- Often, it is desirable to apply different gain and bias values to each channel of a color image separately
 - Example: A color image that utilizes the HSB (Hue-Saturation-Brightness) color model.



Credit “Learn Ui Design Blog”

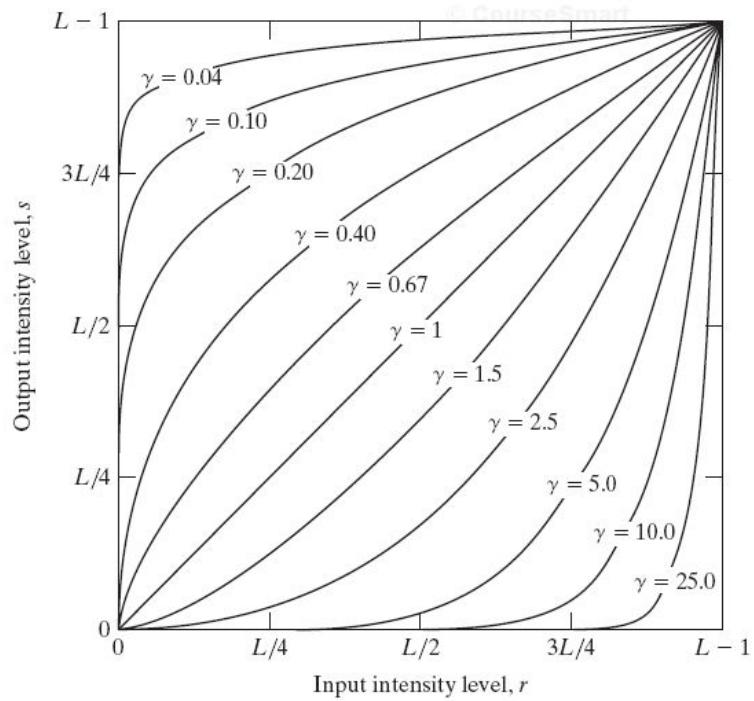
- Since all color information is contained in the H and S channels, it may be useful to adjust ONLY the brightness, encoded in channel B, without altering the color of the image in any way.
- Rescaling the channels of a color image in a non-uniform manner is also possible rescaling each color channel separately.

Example: Gamma Correction

$$S = r^\gamma$$

a
b
c
d

FIGURE 3.9
(a) Aerial image.
(b)–(d) Results of applying the transformation in Eq. (3.2-3) with $c = 1$ and $\gamma = 3.0, 4.0$, and 5.0 , respectively. (Original image for this example courtesy of NASA.)



© CourseSmart

Putting it all together: Gain, Bias, and Gamma

- $C_{\text{out}} = (\alpha C_{\text{in}} + \beta)^\gamma$
- α is known as **gain** (exposure)
- β is known as **bias** (offset)
- γ maps to a non-linear curve (**gamma** correction)

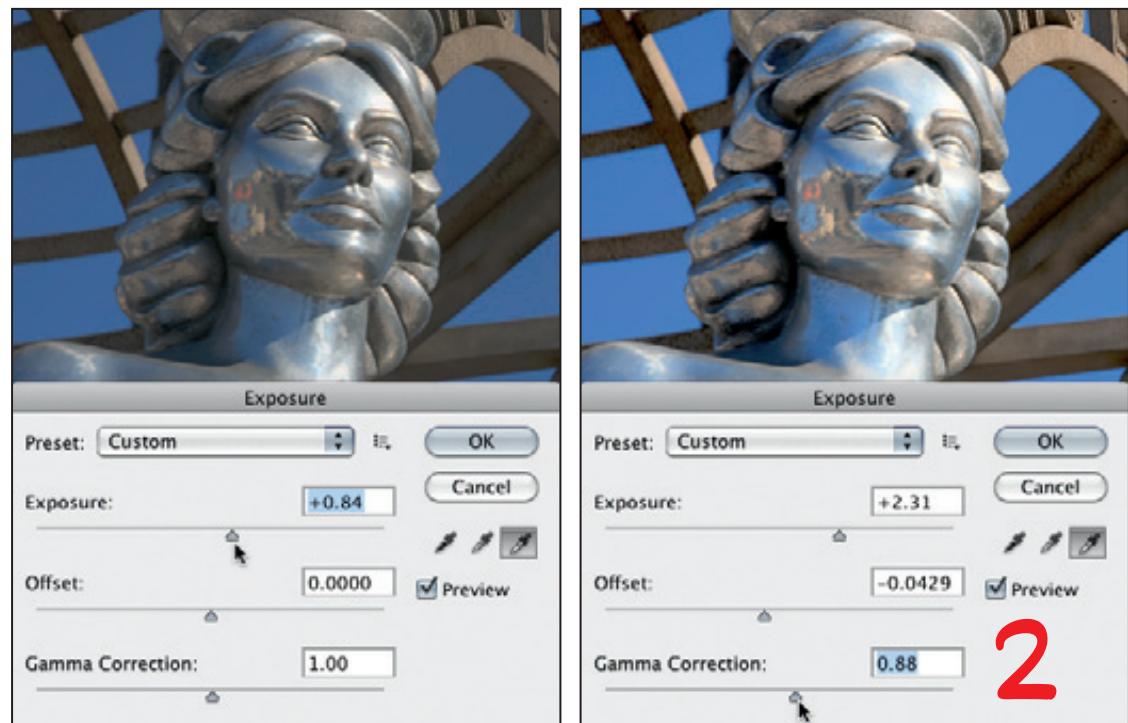


Image of Photoshop from
Christian Bloch - The HDRI Handbook 2.0

Dynamic Range

The World is a High Dynamic Range (HDR)



1:1



1:1,500



1:25,000



1:400,000



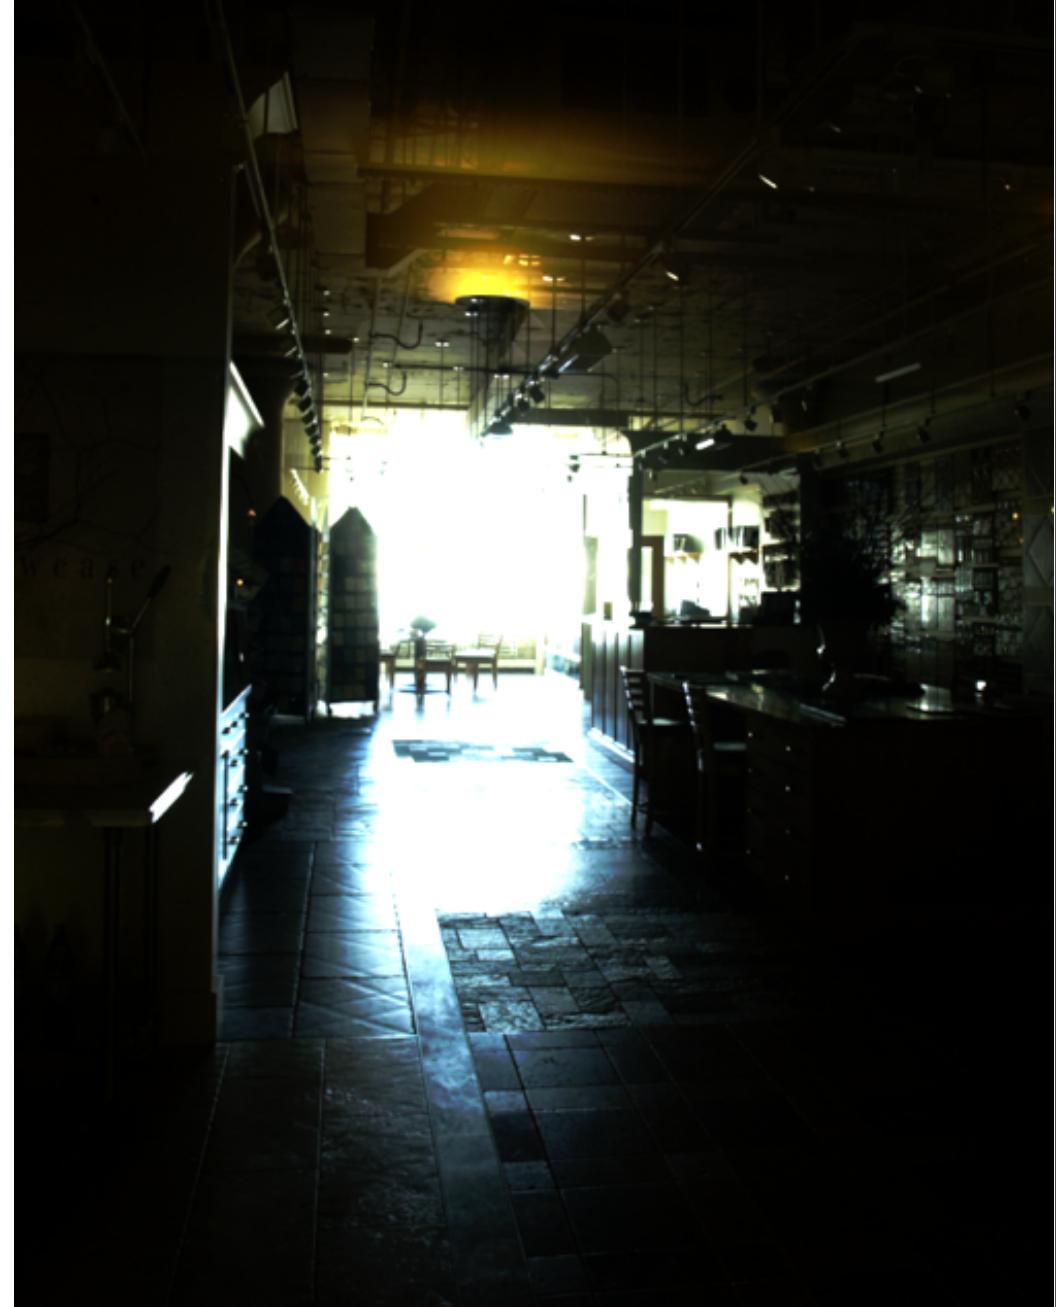
1:2,000,000,000

Examples

**Sun overexposed
Foreground too dark**

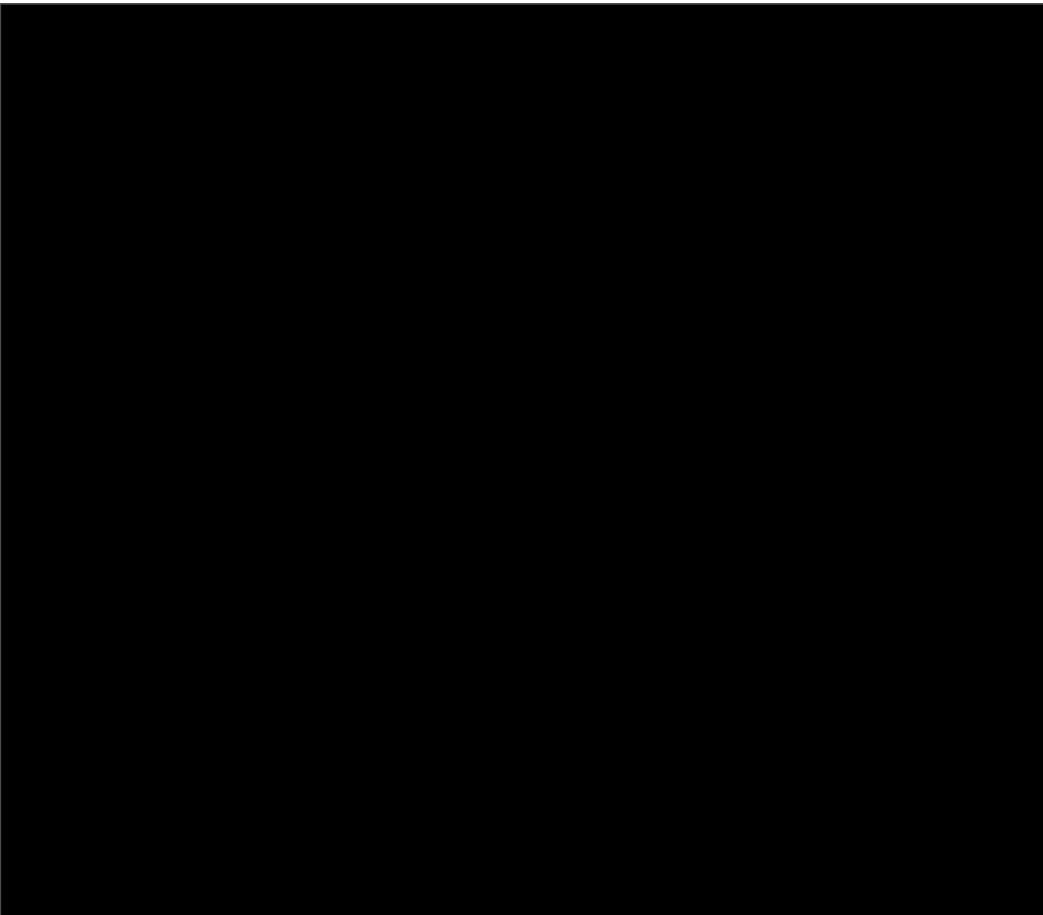


**Inside is too dark
Outside is too bright**



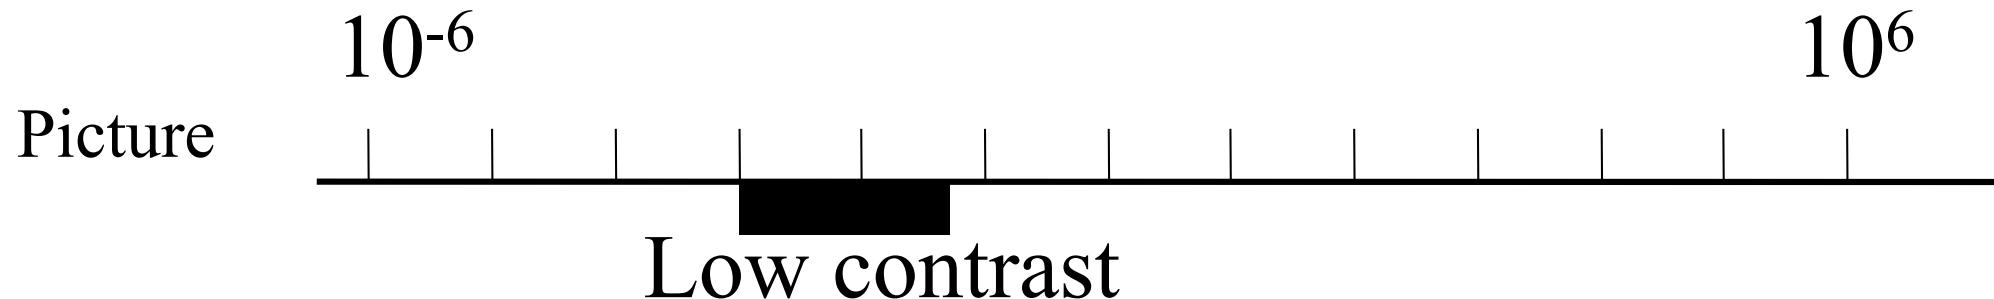
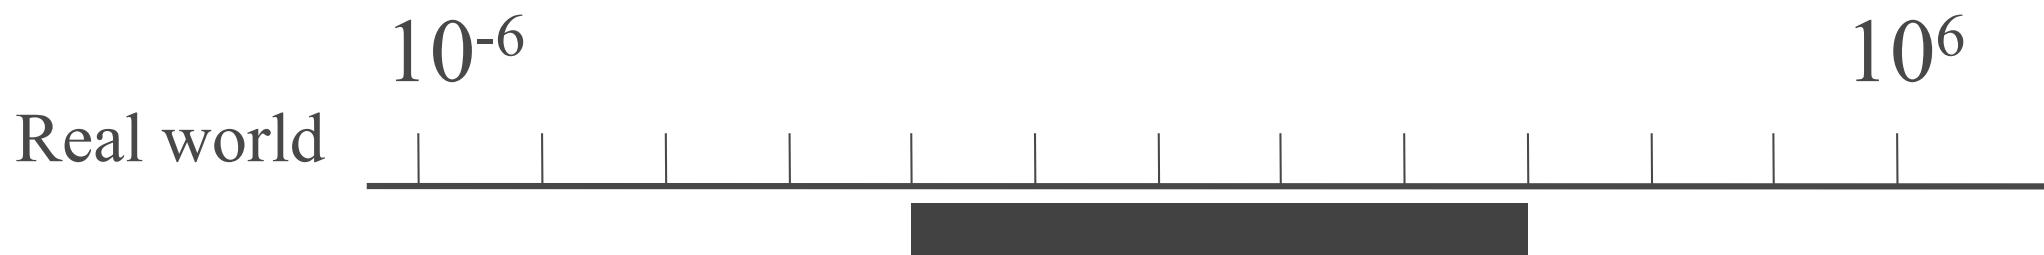
Dynamic Range in Displays?

- Range of pure black vs. pure white?



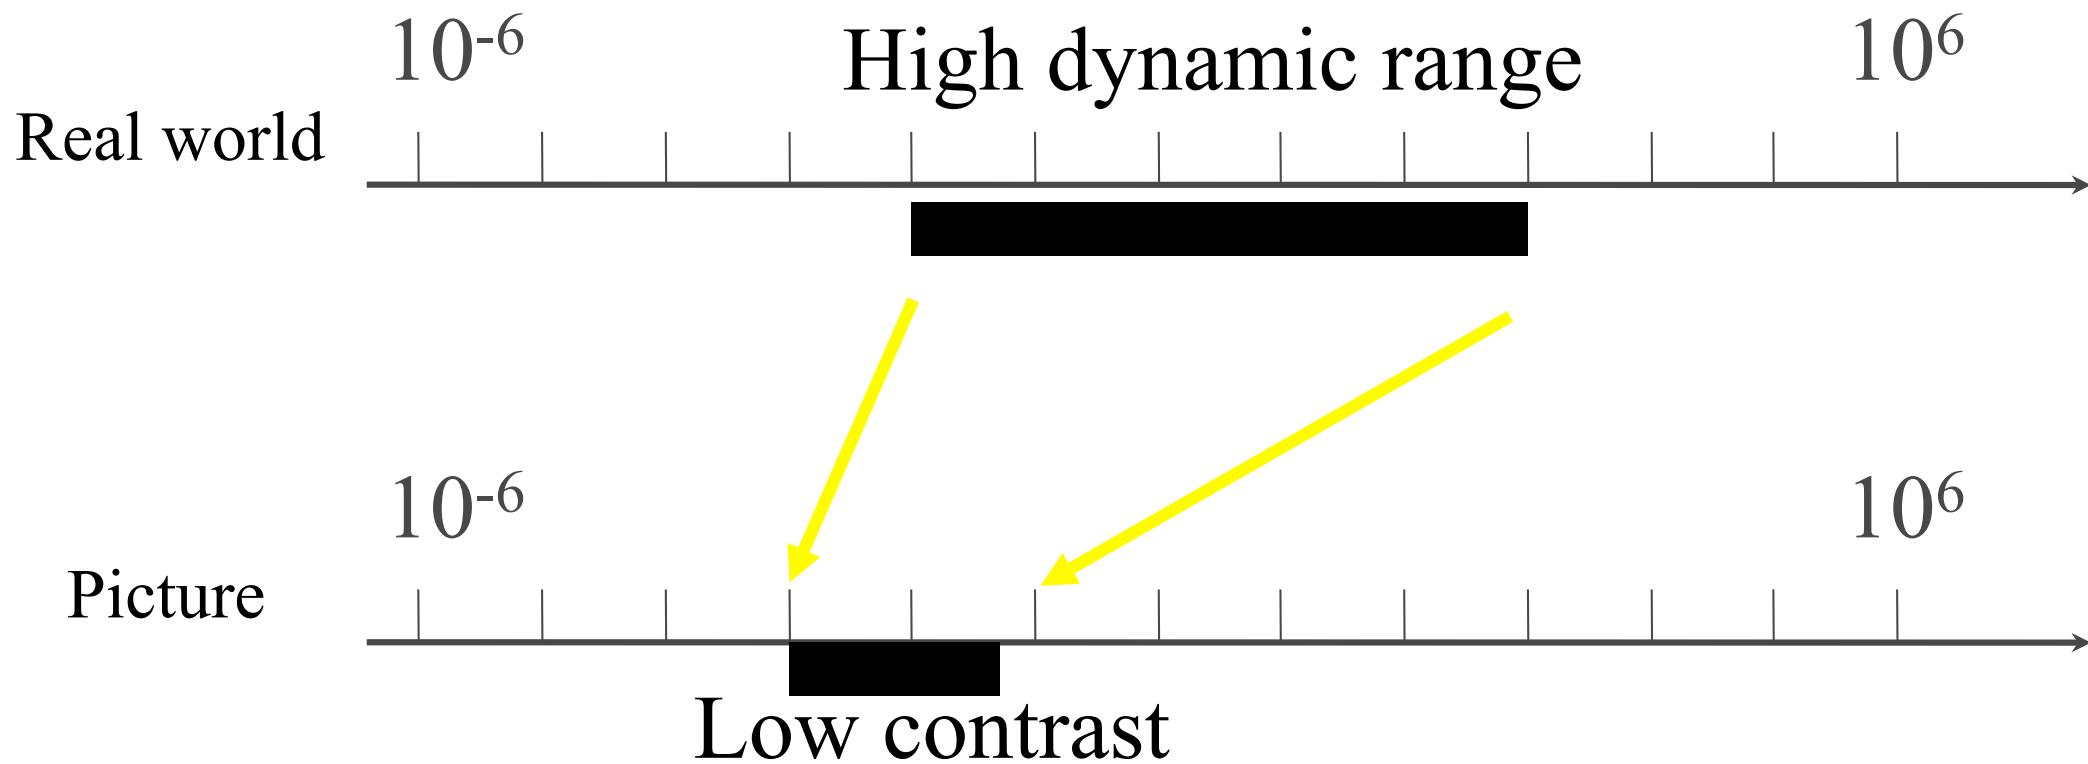
Dynamic Range in Displays?

- Typically 1: 20 or 1:50
 - Black  is ~ 50x darker than white



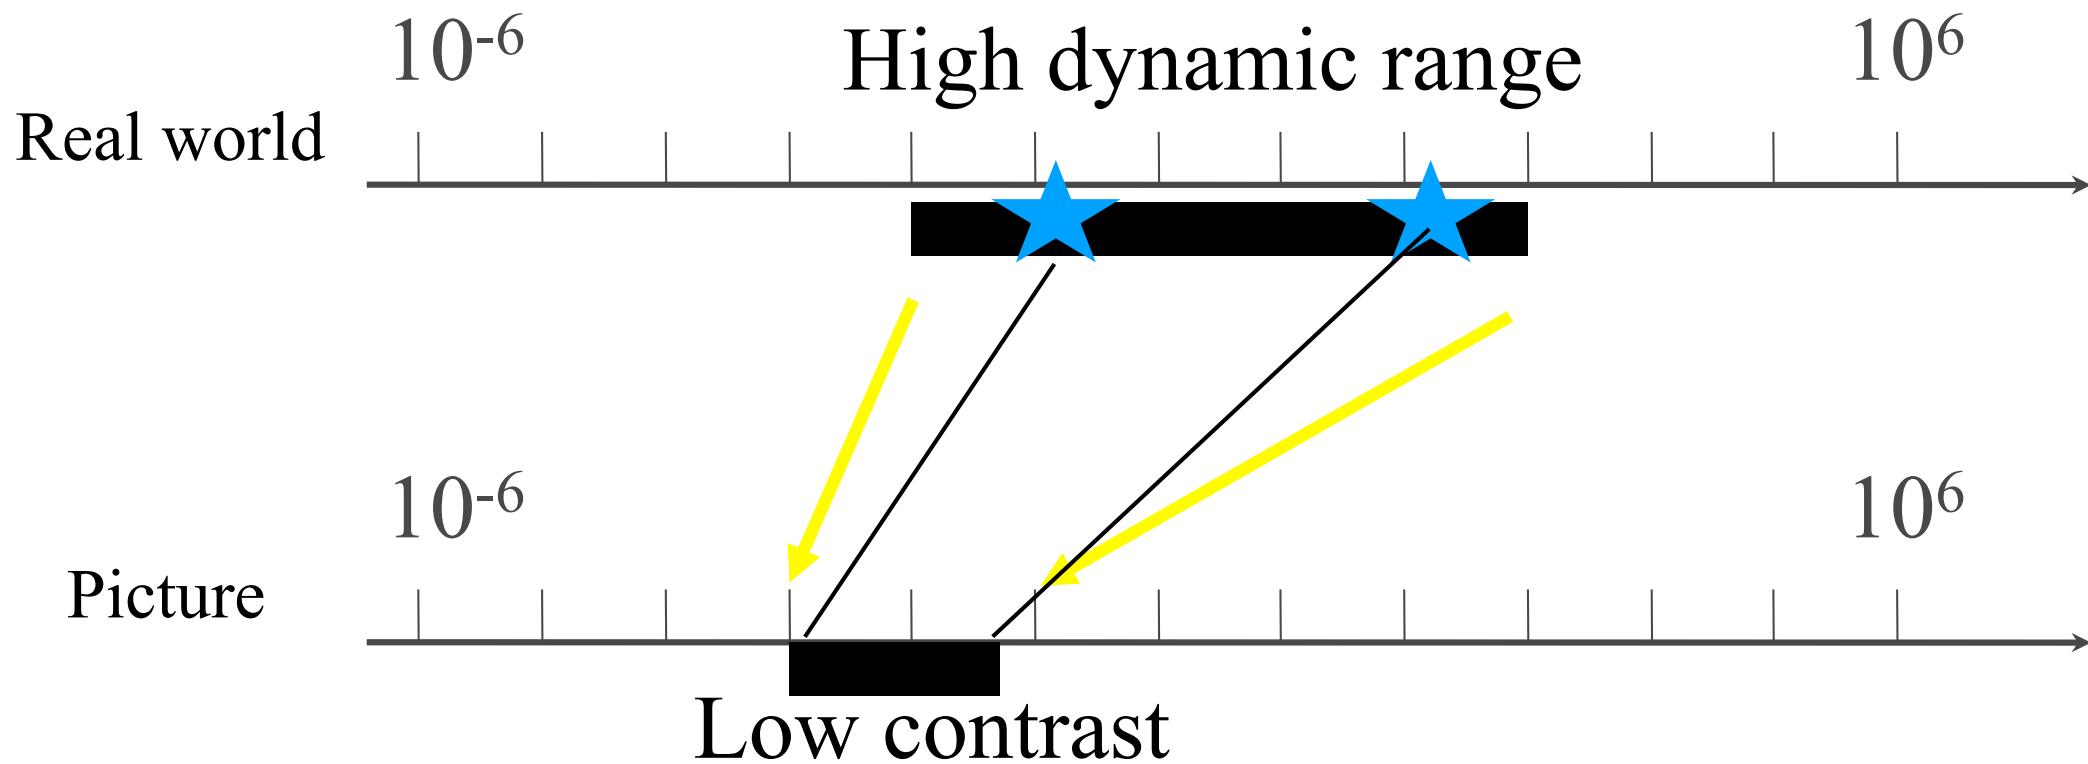
Problem: Displaying the Information

- Problem: How should we map scene radiances (up to 1:100,000) to display radiances (only around 1:100) to produce a satisfactory image?
- Goal: match limited contrast of the display medium while preserving details
- Solution: **Tone Mapping**



First solution: Linear mapping

- We will find the pixels with min and max intensity in the input image.
- Map them to the min and max intensities of the display
- Everything in between is mapped linearly.



Without HDR + Tone Mapping





From Durand and Dorsey. No single global exposure can preserve both the colors of the sky and the details of the landscape, as shown on the rightmost images.

With HDR + Tone Mapping



Before



After



<http://abduzeedo.com/20-beautiful-hdr-pictures>

Not All Tone Mapping Produces Extreme Results, Sometimes Just Beautiful Ones



Check (recommended)

<http://luminous-landscape.com/essays/hdr-plea.shtml>

Tone Mapping

Question: But why do we need more than 100 levels of intensity (luminance) if in the input file we only have 256 values of intensities (RGB) ?

Answer: Not all file format has so few levels.

Even PPM could have 2 bytes per channel, so $256^2=65536$ levels per channel.

Other formats gives much wider range:

Radiance RGBE Format (.hdr)

32 bits/pixel



Red

Green

Blue

Exponent

$$(145, 215, 87, 149) =$$

$$(145, 215, 87) * 2^{(149-128)} =$$

1190000 1760000 713000

$$(145, 215, 87, 103) =$$

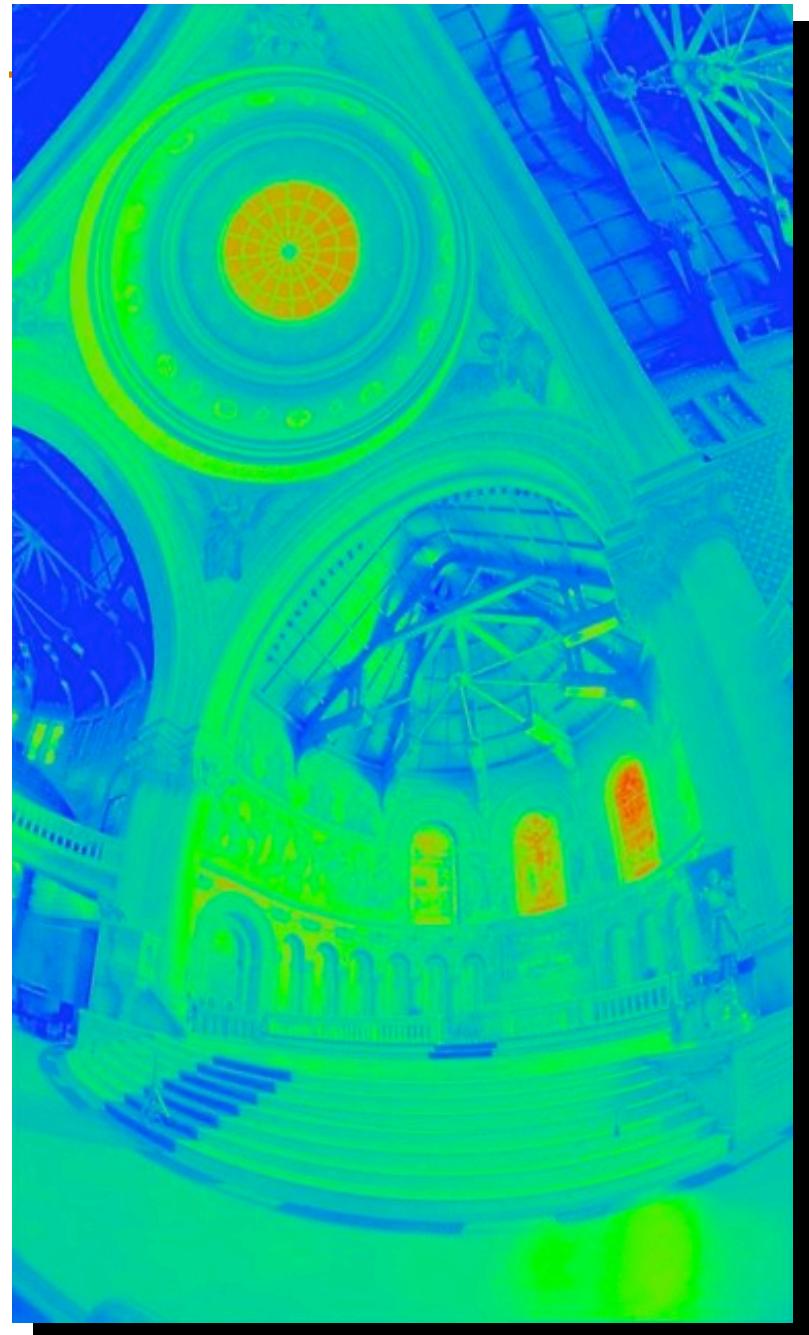
$$(145, 215, 87) * 2^{(103-128)} =$$

0.00000432 0.00000641 0.00000259

Ward, Greg. "Real Pixels," in Graphics Gems IV, edited by James Arvo, Academic Press, 1994

The Radiance Map

Radiance Definition: much of the power (in watts) is emitted by a cm^2 surface will be received by an optical system looking at that surface from a specified angle of view.

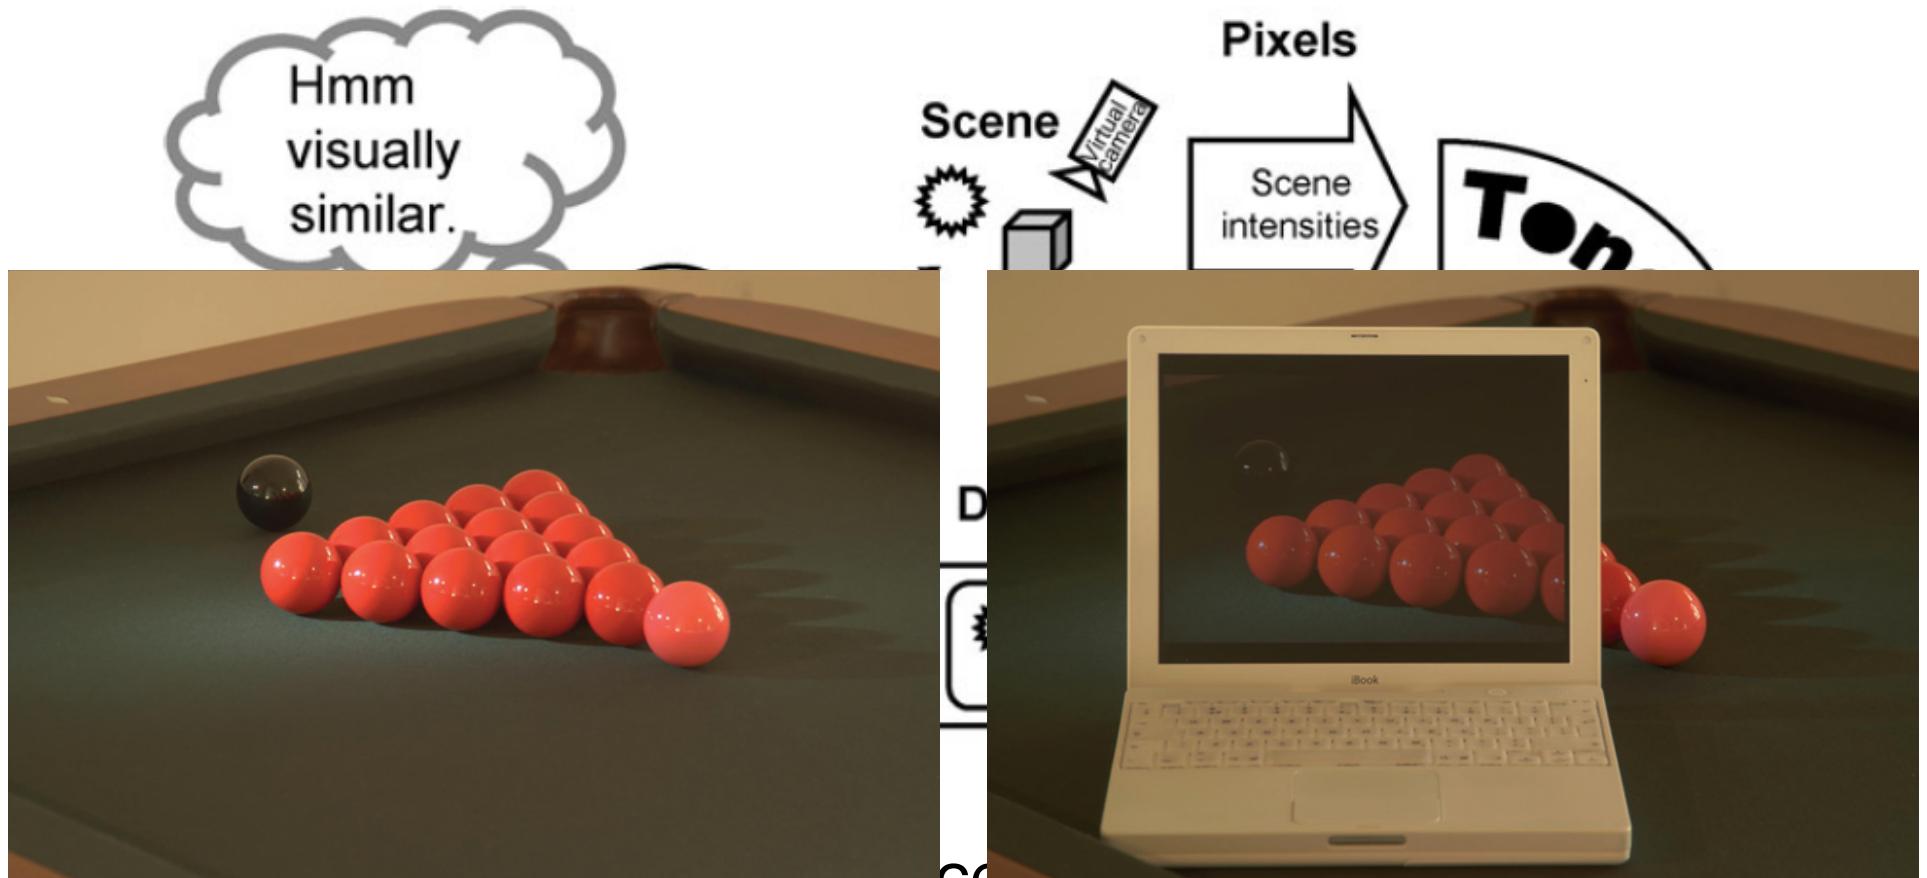


The Radiance Map



Linearly scaled to
display device

Approach: Visual Matching



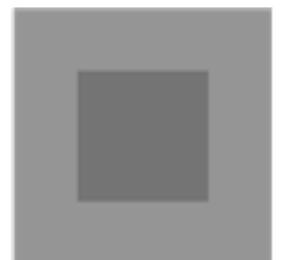
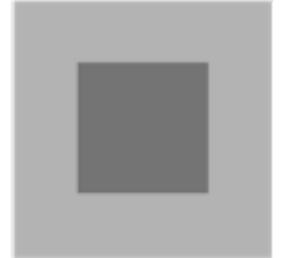
We do not need to reproduce the true radiance as long as it gives us a visual match.

Eyes and Dynamic Range

- We're sensitive to change (multiplicative)
 - A ratio of 1:2 is perceived as the same contrast as a ratio of 100 to 200
 - Use the log domain as much as possible
- But, eyes are **not** photometers
 - Dynamic adaptation (very local in retina)
 - Different sensitivity to spatial frequencies



Headlights
are ON in
both
photos



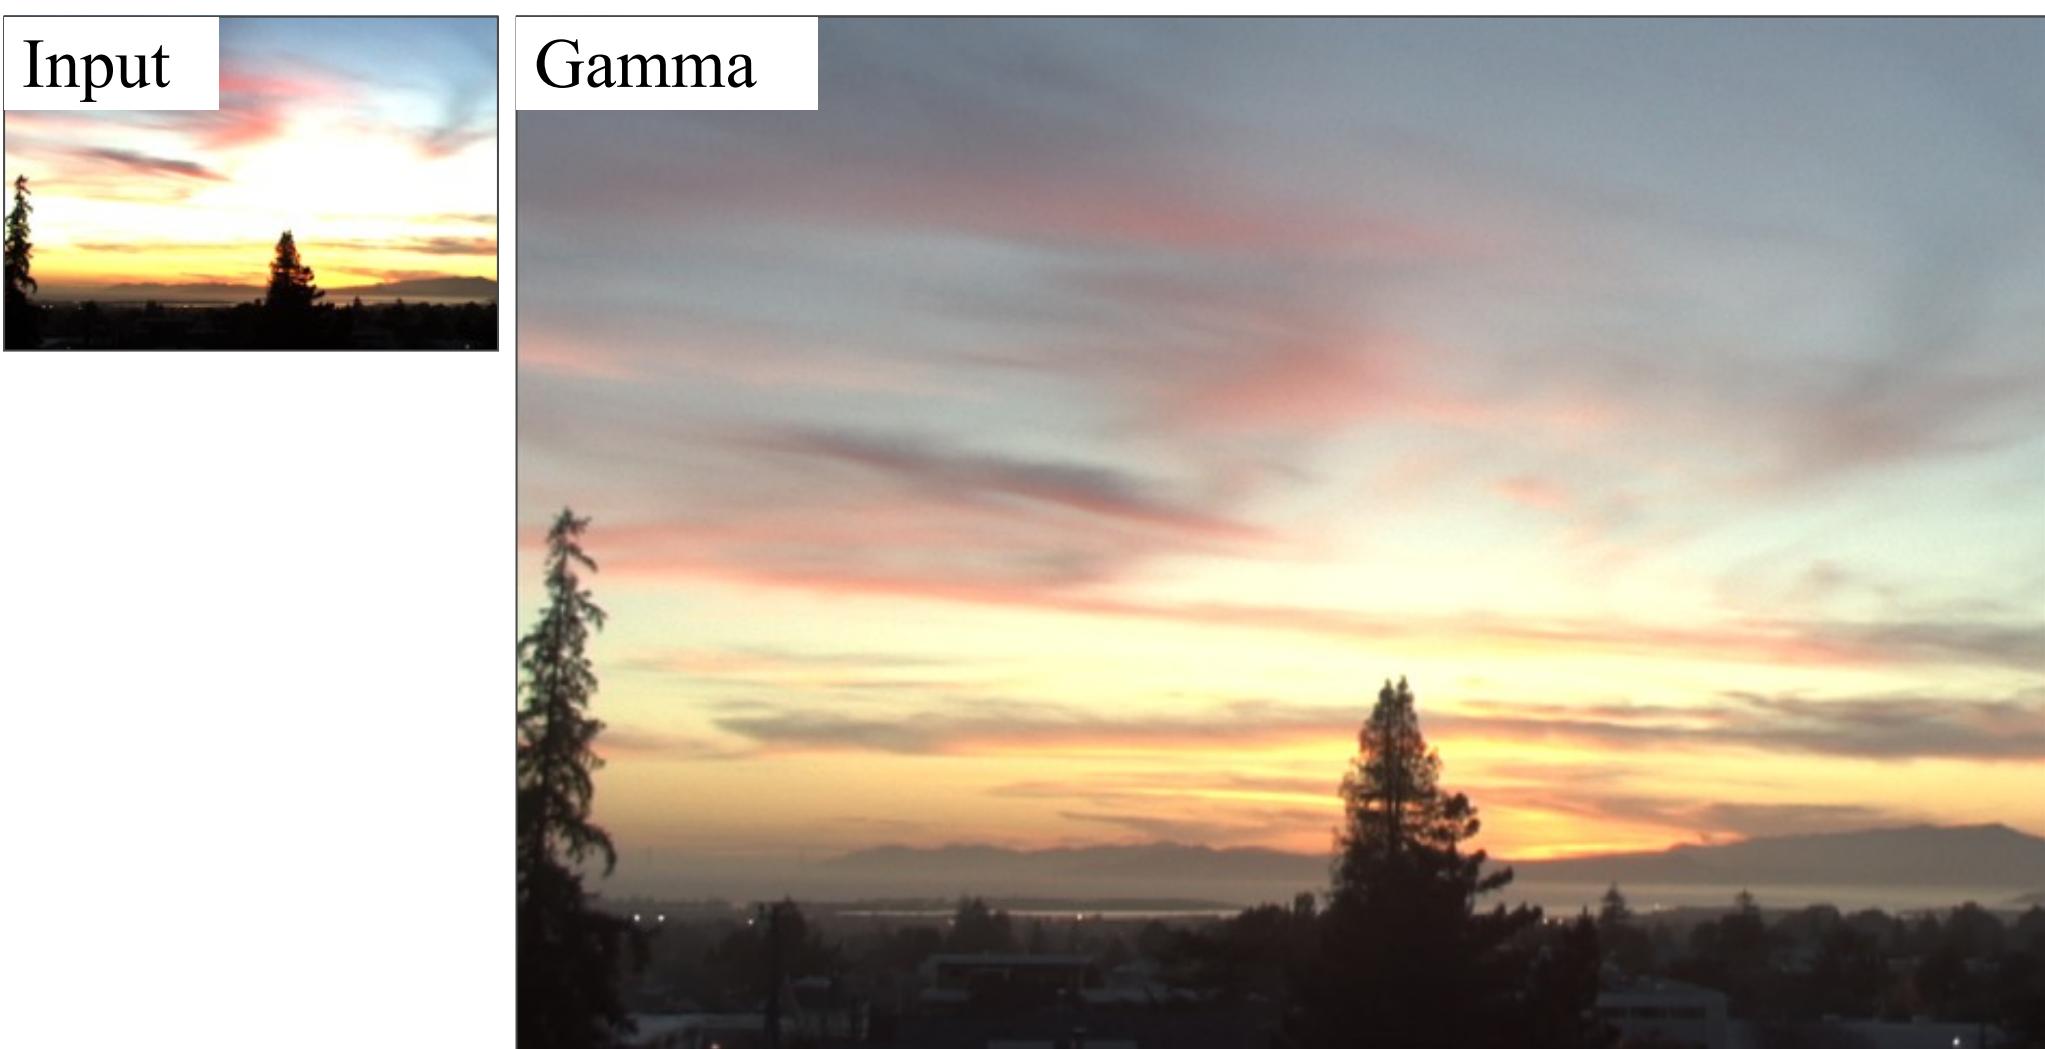
Can we just scale? Maybe!

- For a color image, try to convert the input (world) luminance \mathbf{L}_w to a target display luminance \mathbf{L}_d
- This type of scaling works (sometimes). In particular, it works best in the log and/or exponential domains
- $\log_{10}(x)=1+\log_{10}(y)$ means $x=10y$
- The base of the log is not important, as long as we are consistent in the mapping

$$\begin{bmatrix} R_d \\ G_d \\ B_d \end{bmatrix} = \begin{bmatrix} L_d & \frac{R_w}{L_w} \\ L_d & \frac{G_w}{L_w} \\ L_d & \frac{B_w}{L_w} \end{bmatrix}$$

What scale value to use? How about Gamma compression

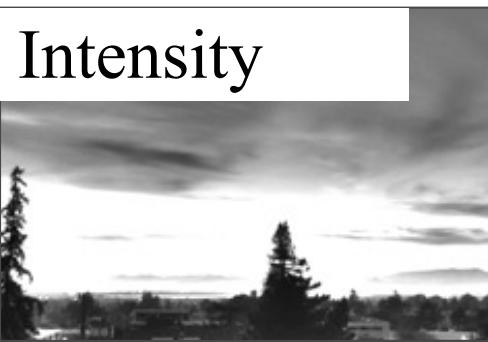
- $C_{\text{out}} = C_{\text{in}}^{\gamma}$, where $0 < \gamma < 1$ applied to each R,G,B channel
- Colors are washed out, why?



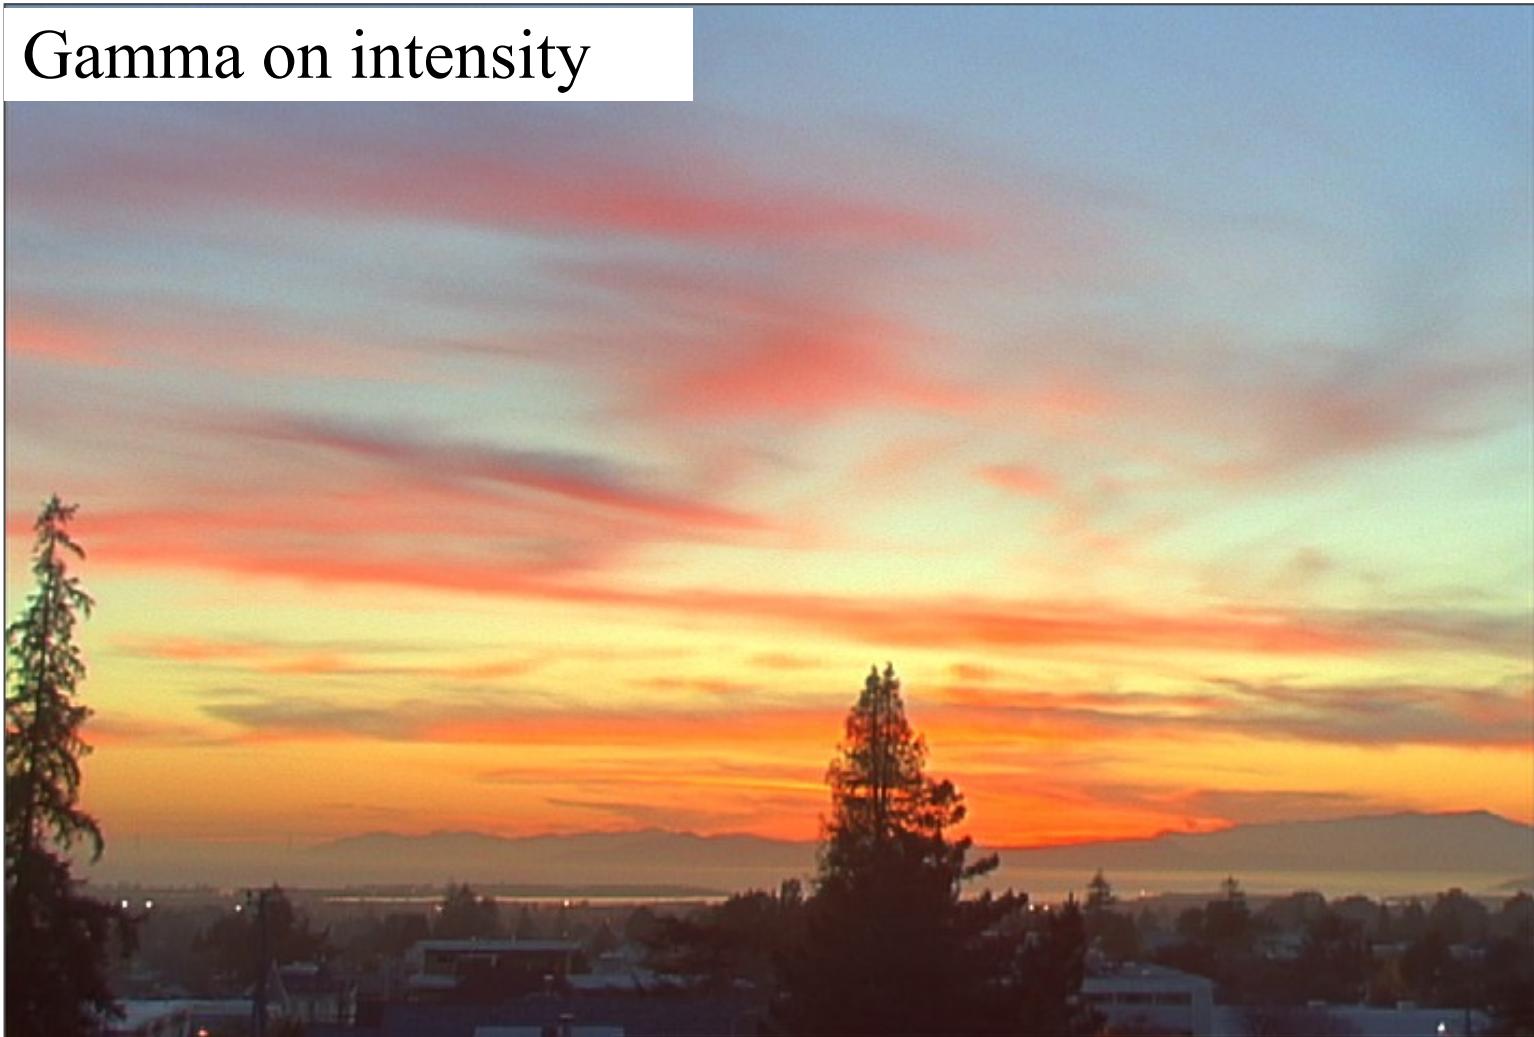
Gamma compression on Intensity

- Colors ok, but details in intensity are blurry

Intensity



Gamma on intensity



Color



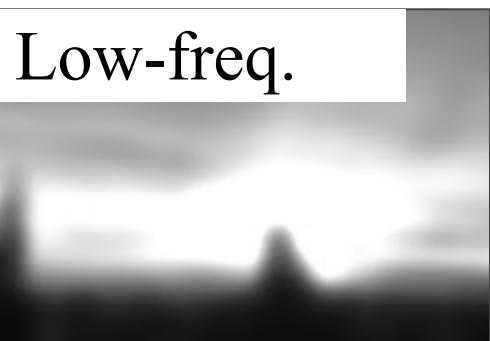
Bilateral Filtering

(slides from Frédo Durand)

Oppenheim 1968, Chiu et al. 1993

- Reduce contrast of low-frequencies
- Keep mid and high frequencies

Low-freq.



Reduce low frequency



High-freq.



Color

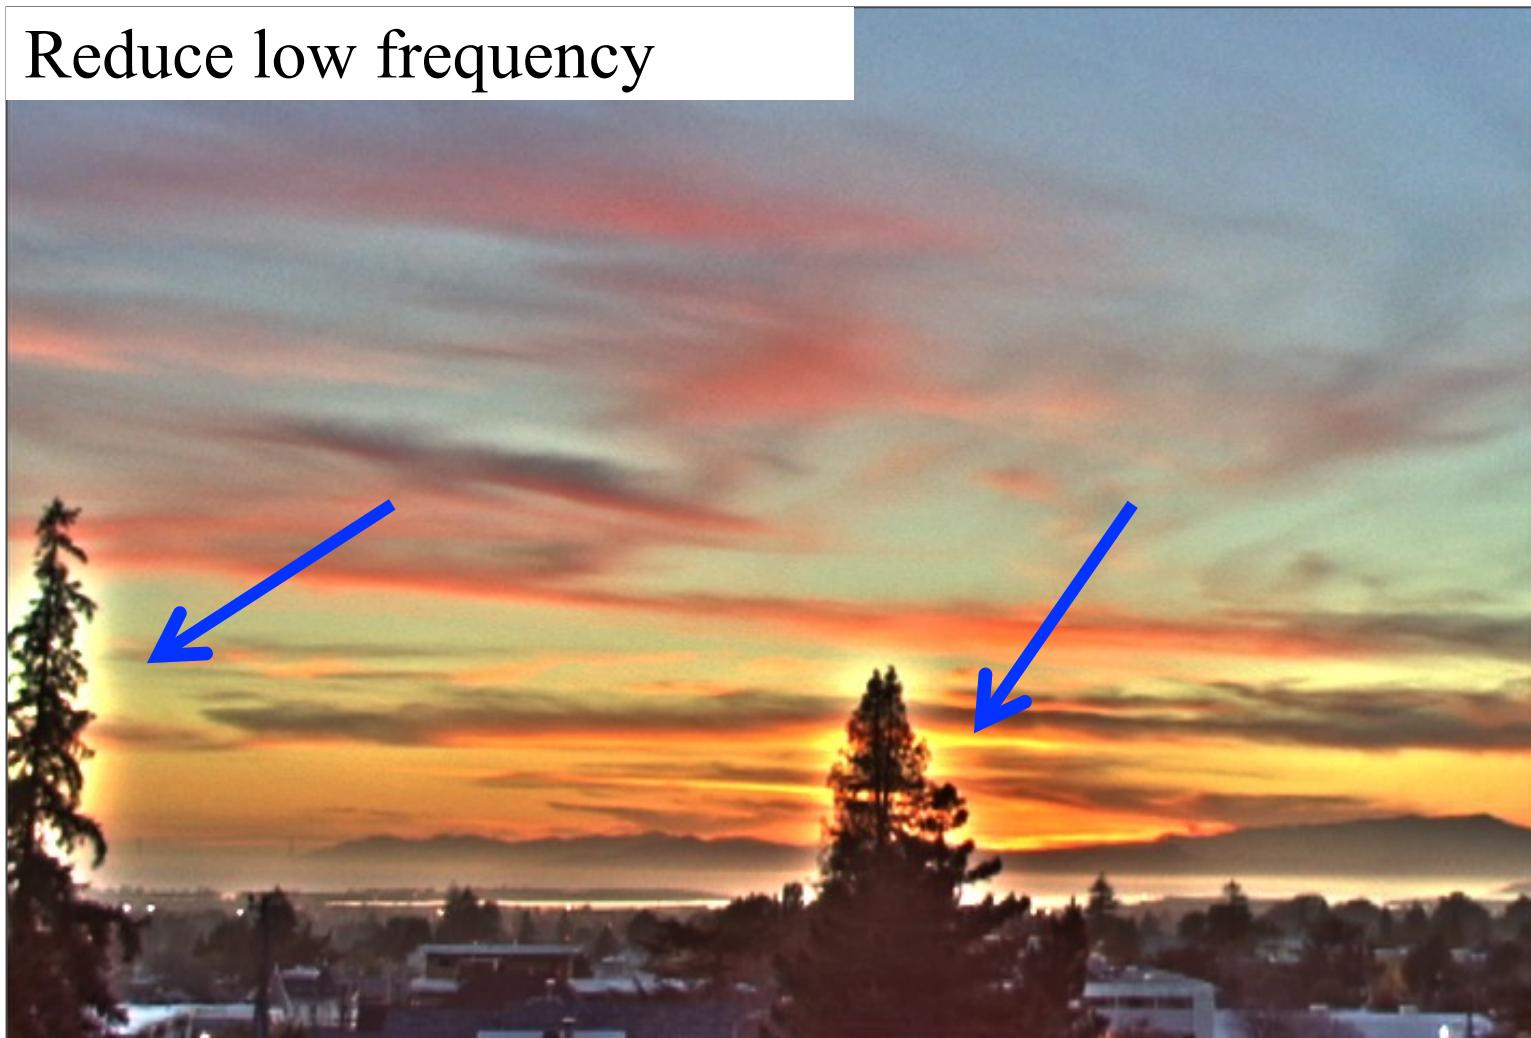


The halo nightmare

- For strong edges
- Because they contain high frequency



Reduce low frequency



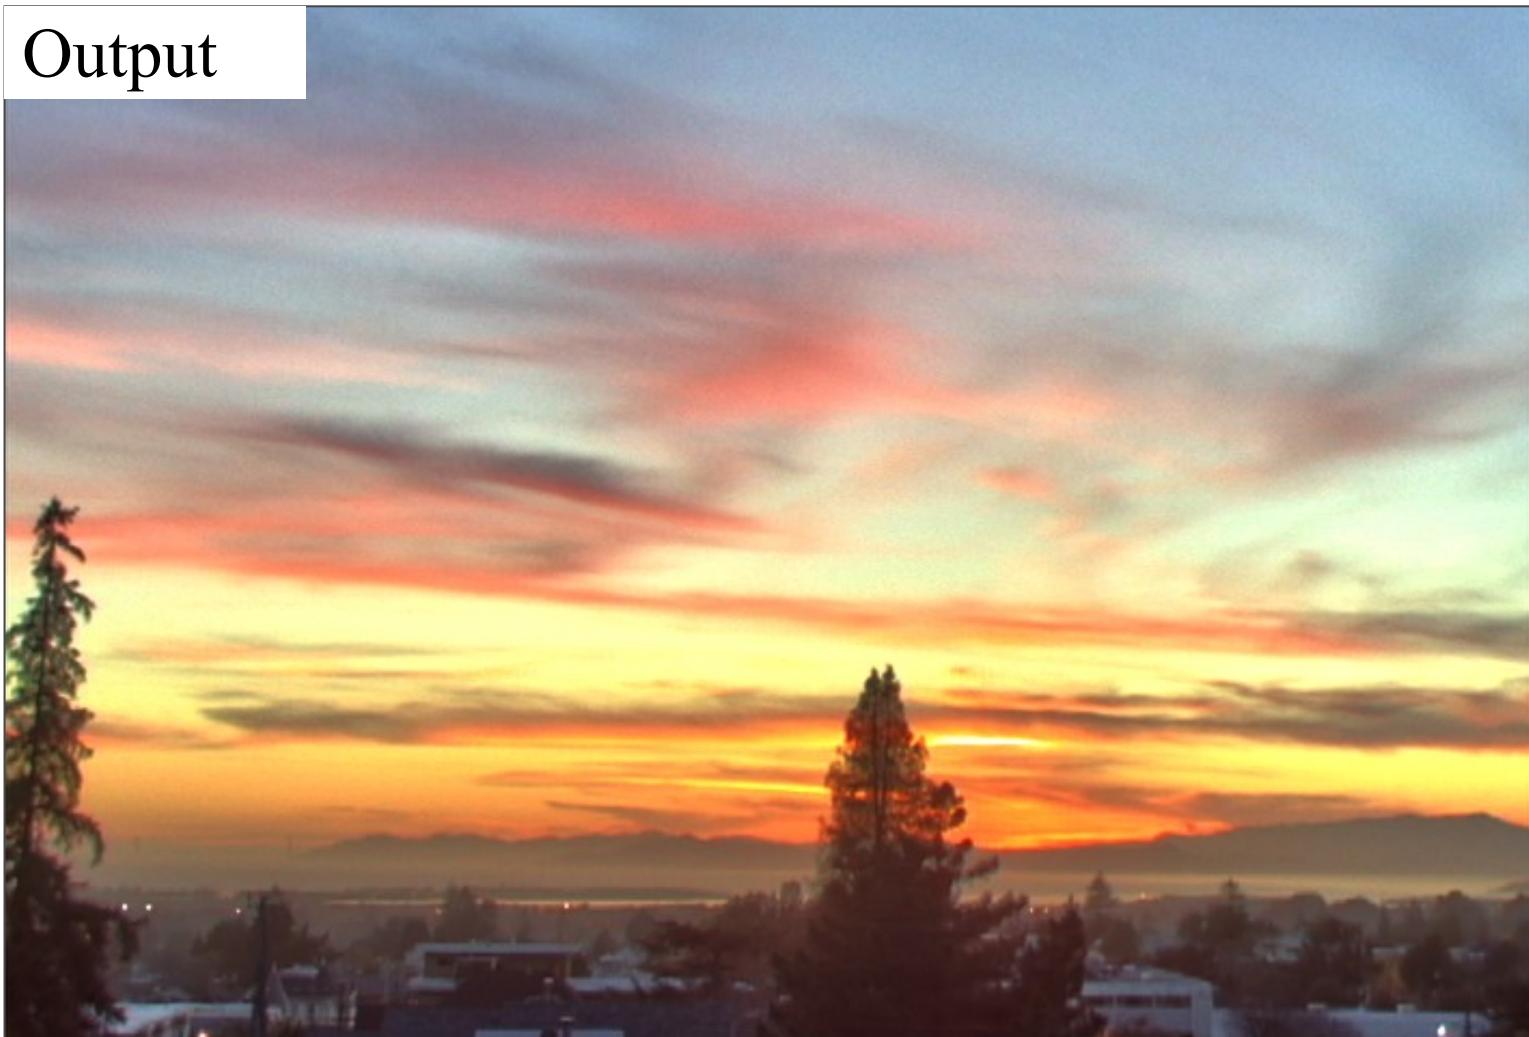
Our approach

- Do not blur across edges
- Non-linear filtering

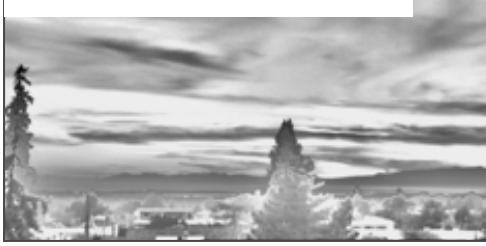
Large-scale



Output



Detail



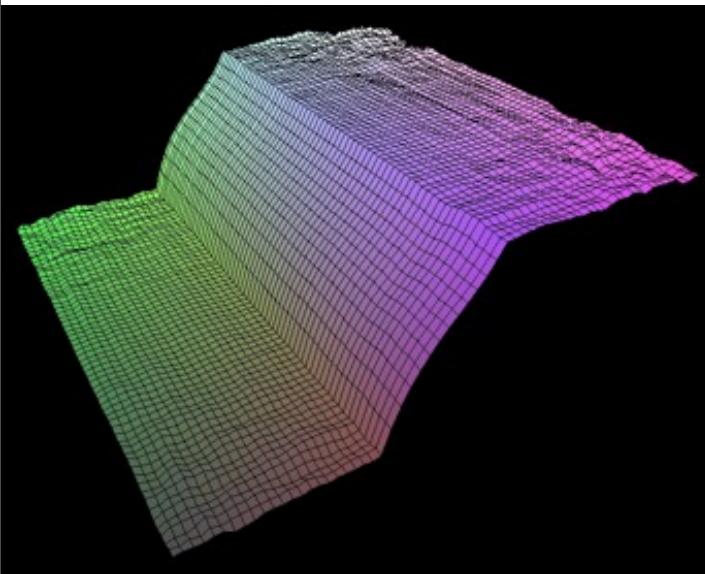
Color



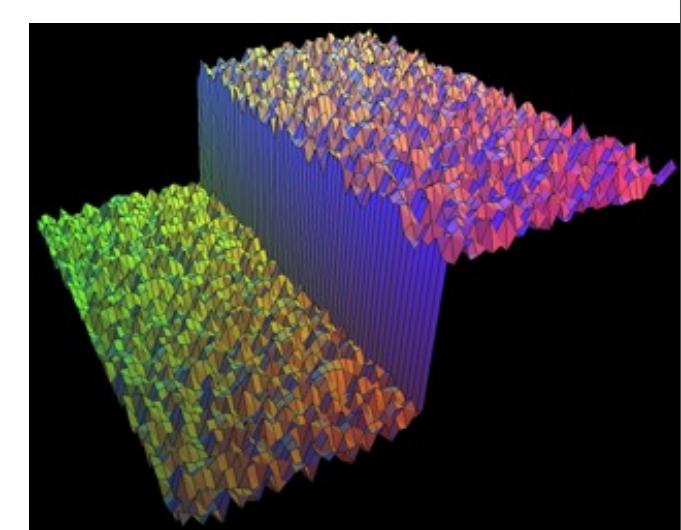
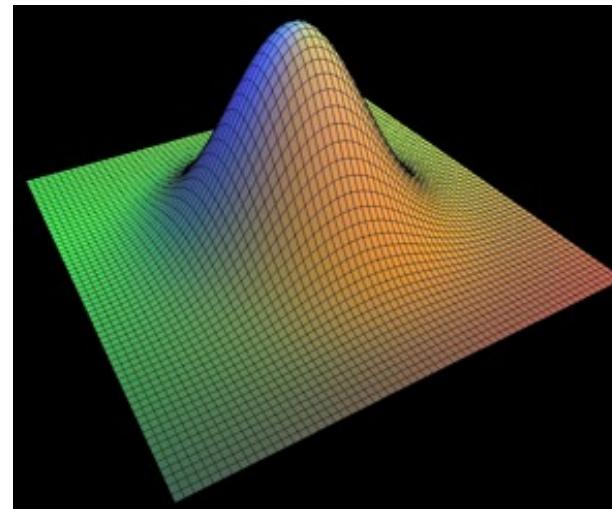
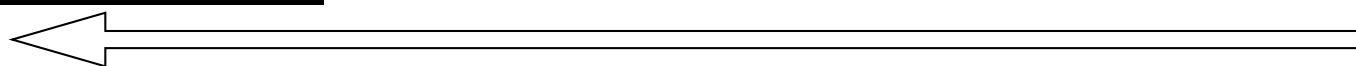
Start with Gaussian filtering

- Here, input is a step function + noise

$$J = f \otimes I$$



output

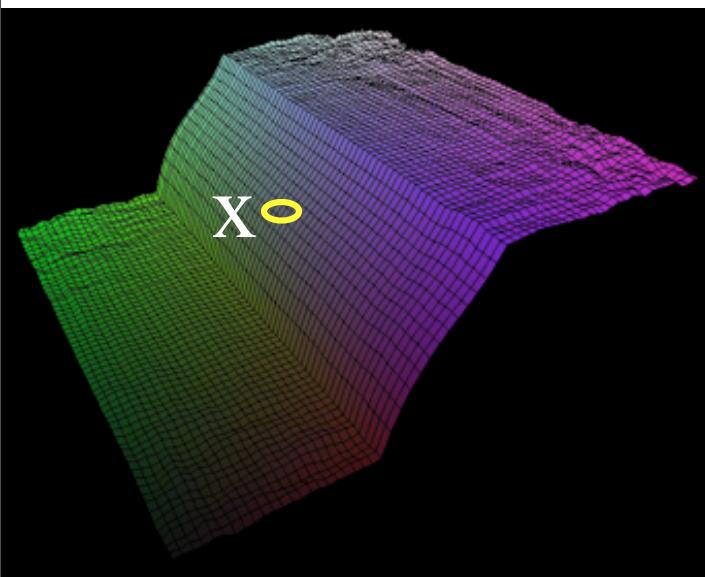


input

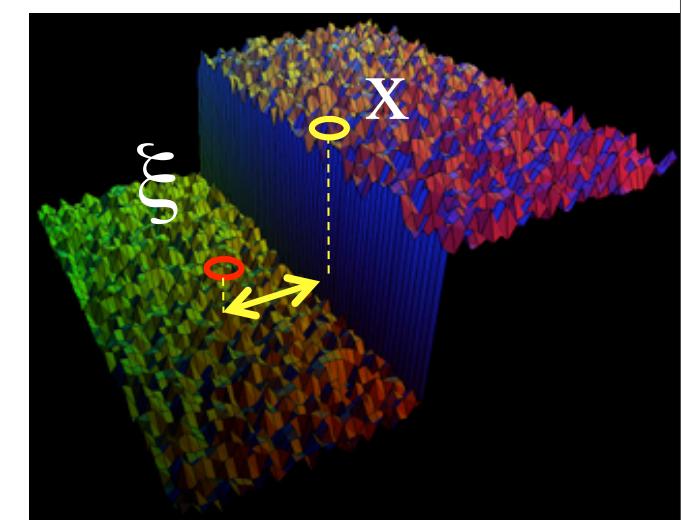
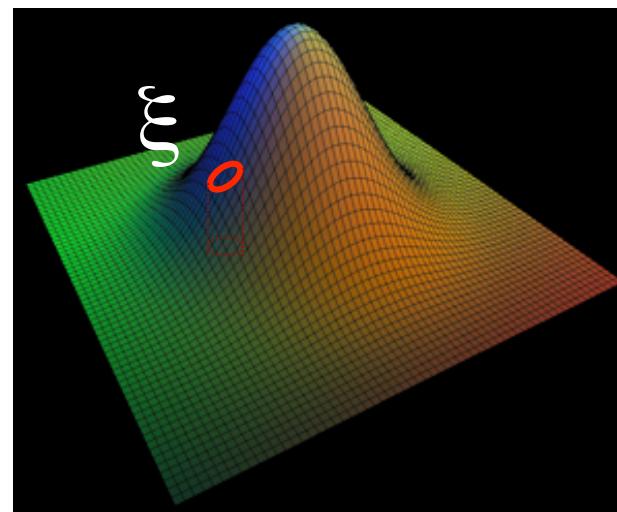
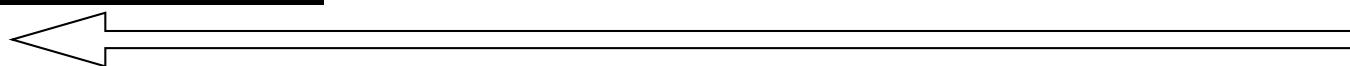
Gaussian filter as weighted average

- Weight of ξ depends on distance to x

$$J(x) = \sum_{\xi} f(x, \xi) I(\xi)$$



output



input

The importance of convex combinations

weight function

$$J(\mathbf{x}) = \sum \underbrace{\widehat{f(\mathbf{x}, \xi)}}_{I(\xi)}$$

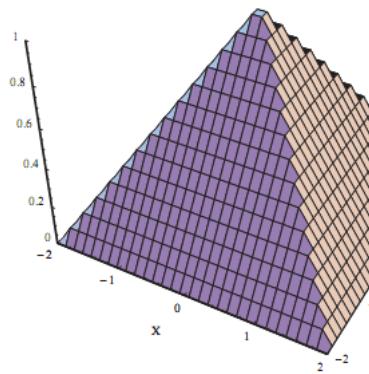
Intensity

x is the point where we need the answer
 ξ is nearby point

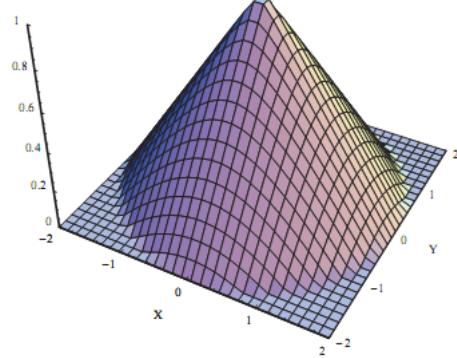
When we smooth, or interpolate we usually use weighted average.

$$G(\mathbf{x}, \xi.x, \xi.y) = \frac{1}{2\pi\sigma} e^{-\frac{1}{2\sigma^2}(x-\xi.x)^2 + (y-\xi.y)^2}$$

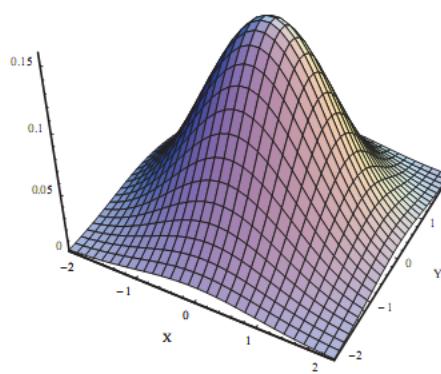
Which functions could $f(\mathbf{x}, \xi)$ be -



(a) Pyramid.



(b) Cone.



(c) Gaussian.

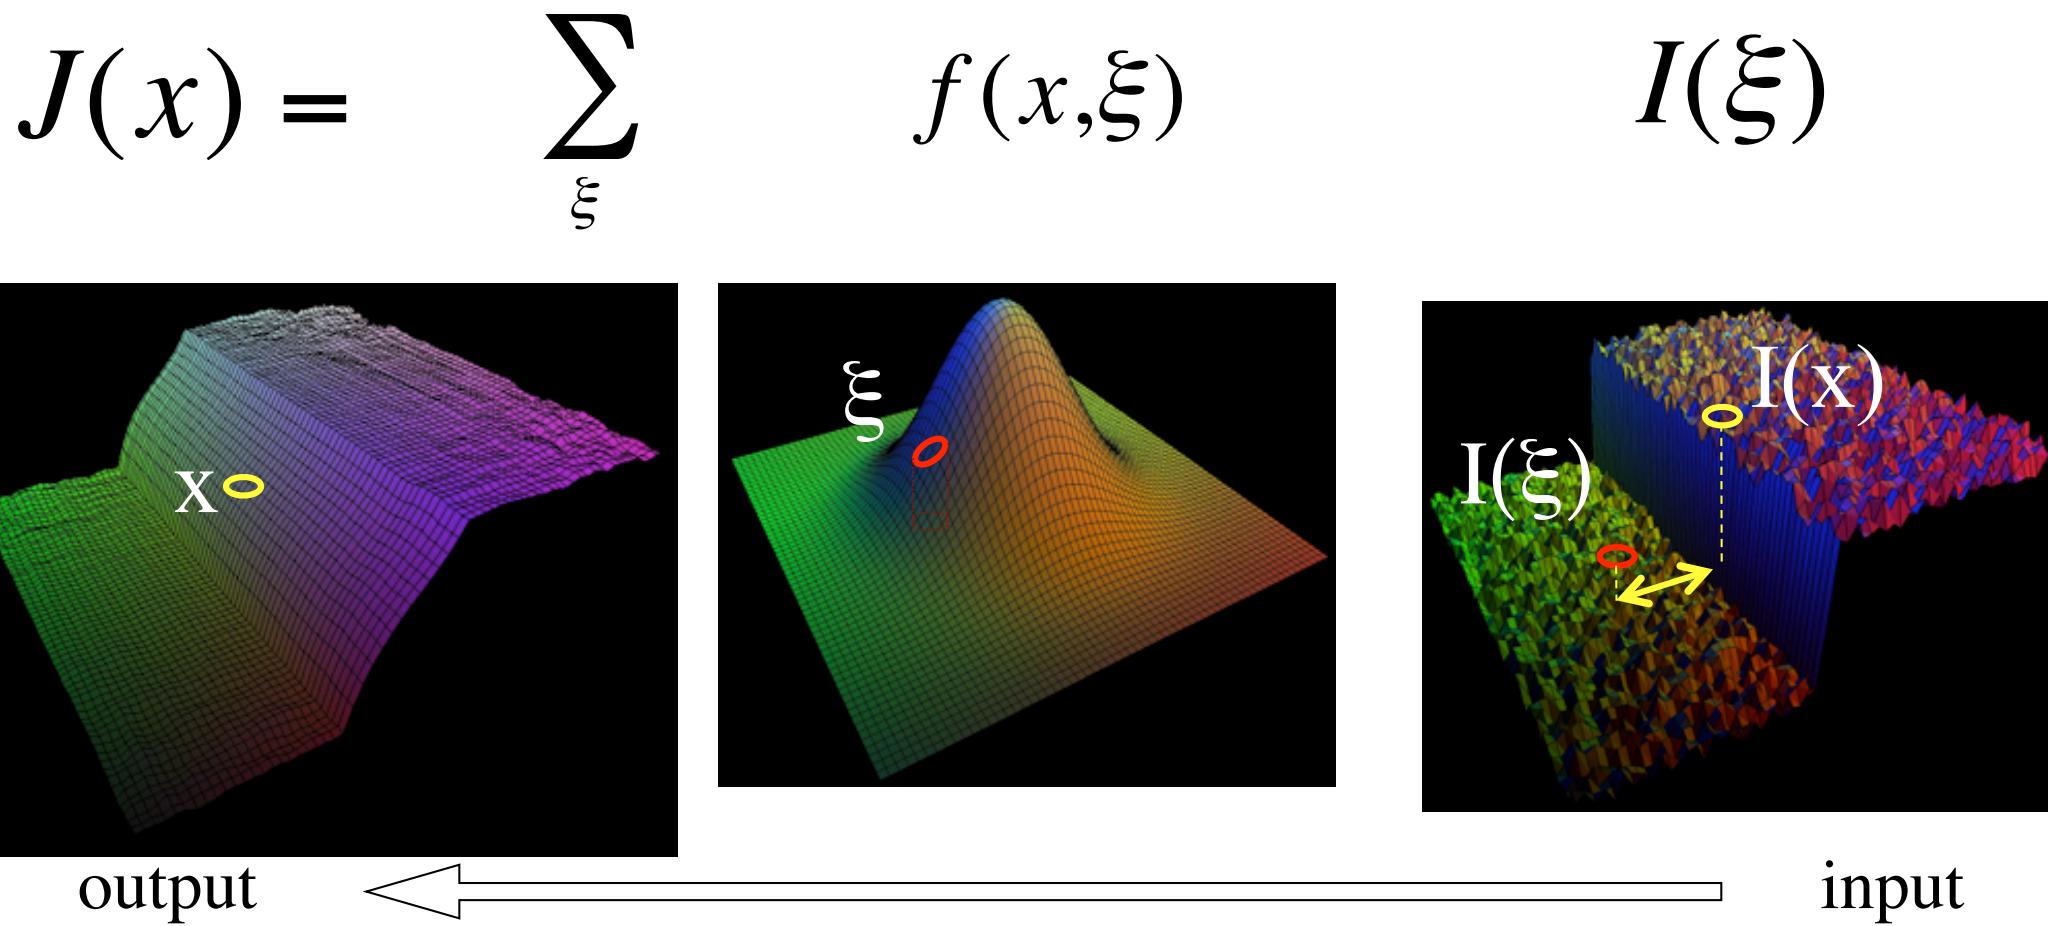
$$f(\mathbf{x}, \xi) = \max \left\{ 0, \frac{3}{\alpha^2} - \frac{3}{\alpha^3} \max(|x - \xi.a|, |y - \xi.b|) \right\}$$

α is the width of the base of the pyramid. So in Fig(a), $\alpha=4$

We will try to make sure that sum of weights = 1 (this is called **convex combination**)
 See example of bilinear interpolation on the whiteboard

The problem of edges

- Here, $I(\xi)$ “pollutes” our estimate $J(x)$
- It is too different



Principle of Bilateral filtering

[Tomasi and Manduchi 1998]

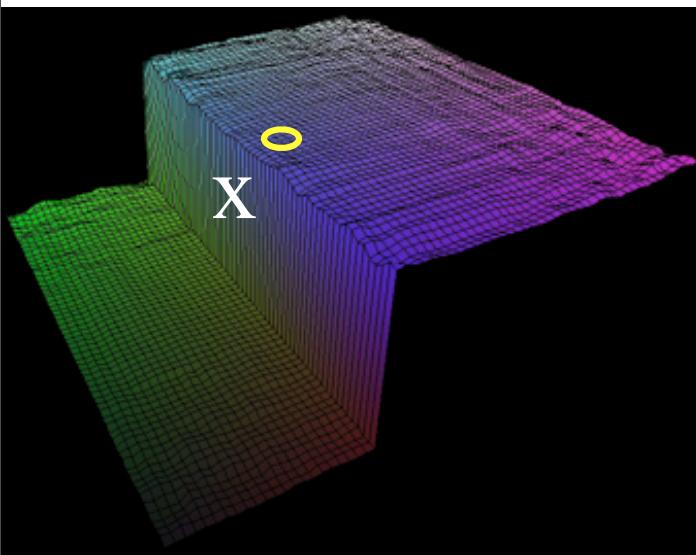
- Penalty **g** on the intensity difference

Remember that the sum of weights must be 1.

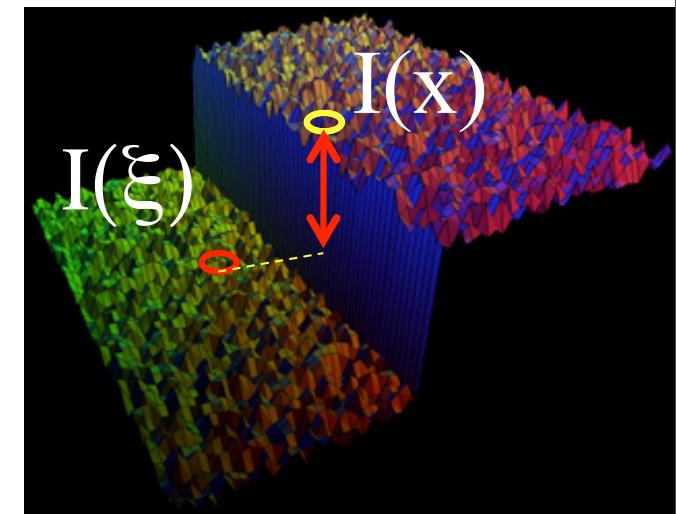
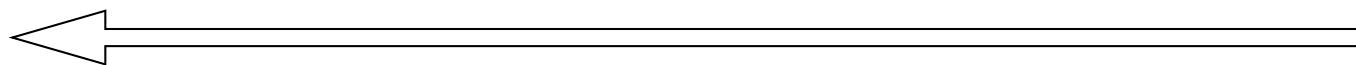
What to do if we skip some terms ?

We will divide the total sum by $k(x)$

$$J(x) = \frac{1}{k(x)} \sum_{\xi} f(x, \xi) \cdot g(I(\xi) - I(x)) \cdot \frac{I(\xi)}{k(x)}$$



output



input

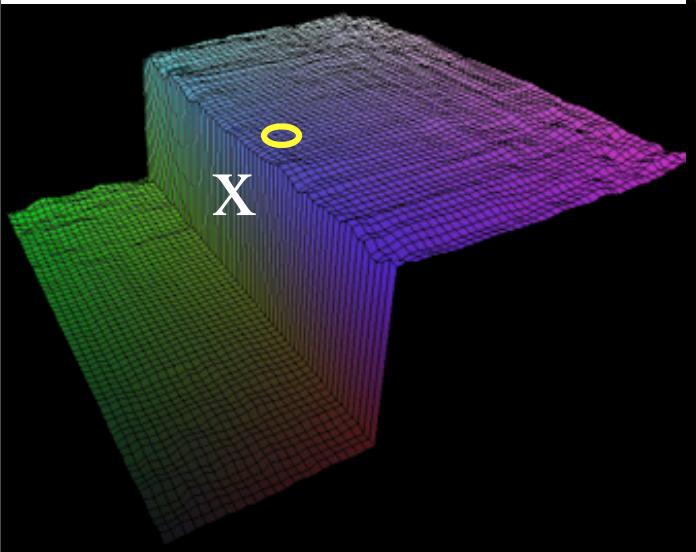
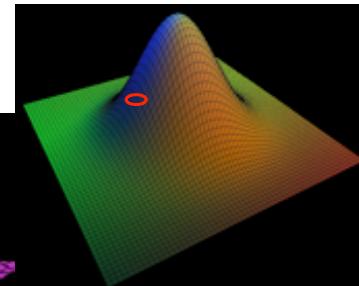
Bilateral filtering

[Tomasi and Manduchi 1998]

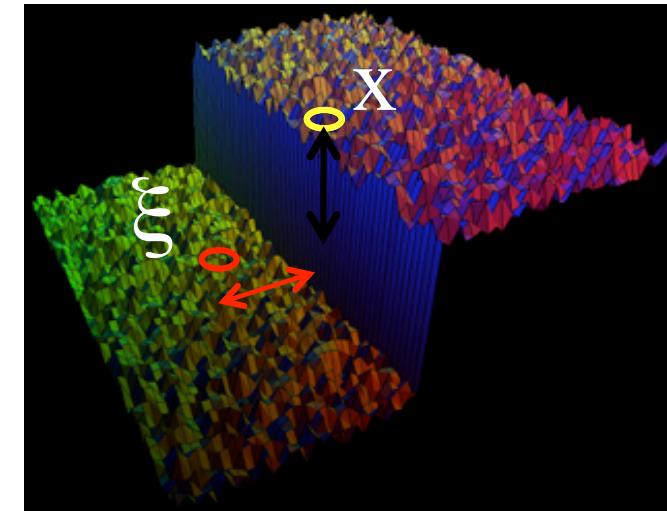
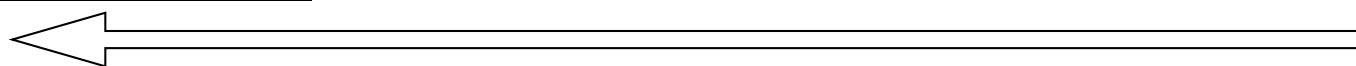
- Spatial Gaussian f

Remember that the sum of weights must be 1.
 What to do if we skip some terms ?
 We will divide the total sum by $k(x)$

$$J(x) = \frac{1}{k(x)} \sum_{\xi} f(x, \xi) g(I(\xi) - I(x)) \quad I(\xi)$$



output



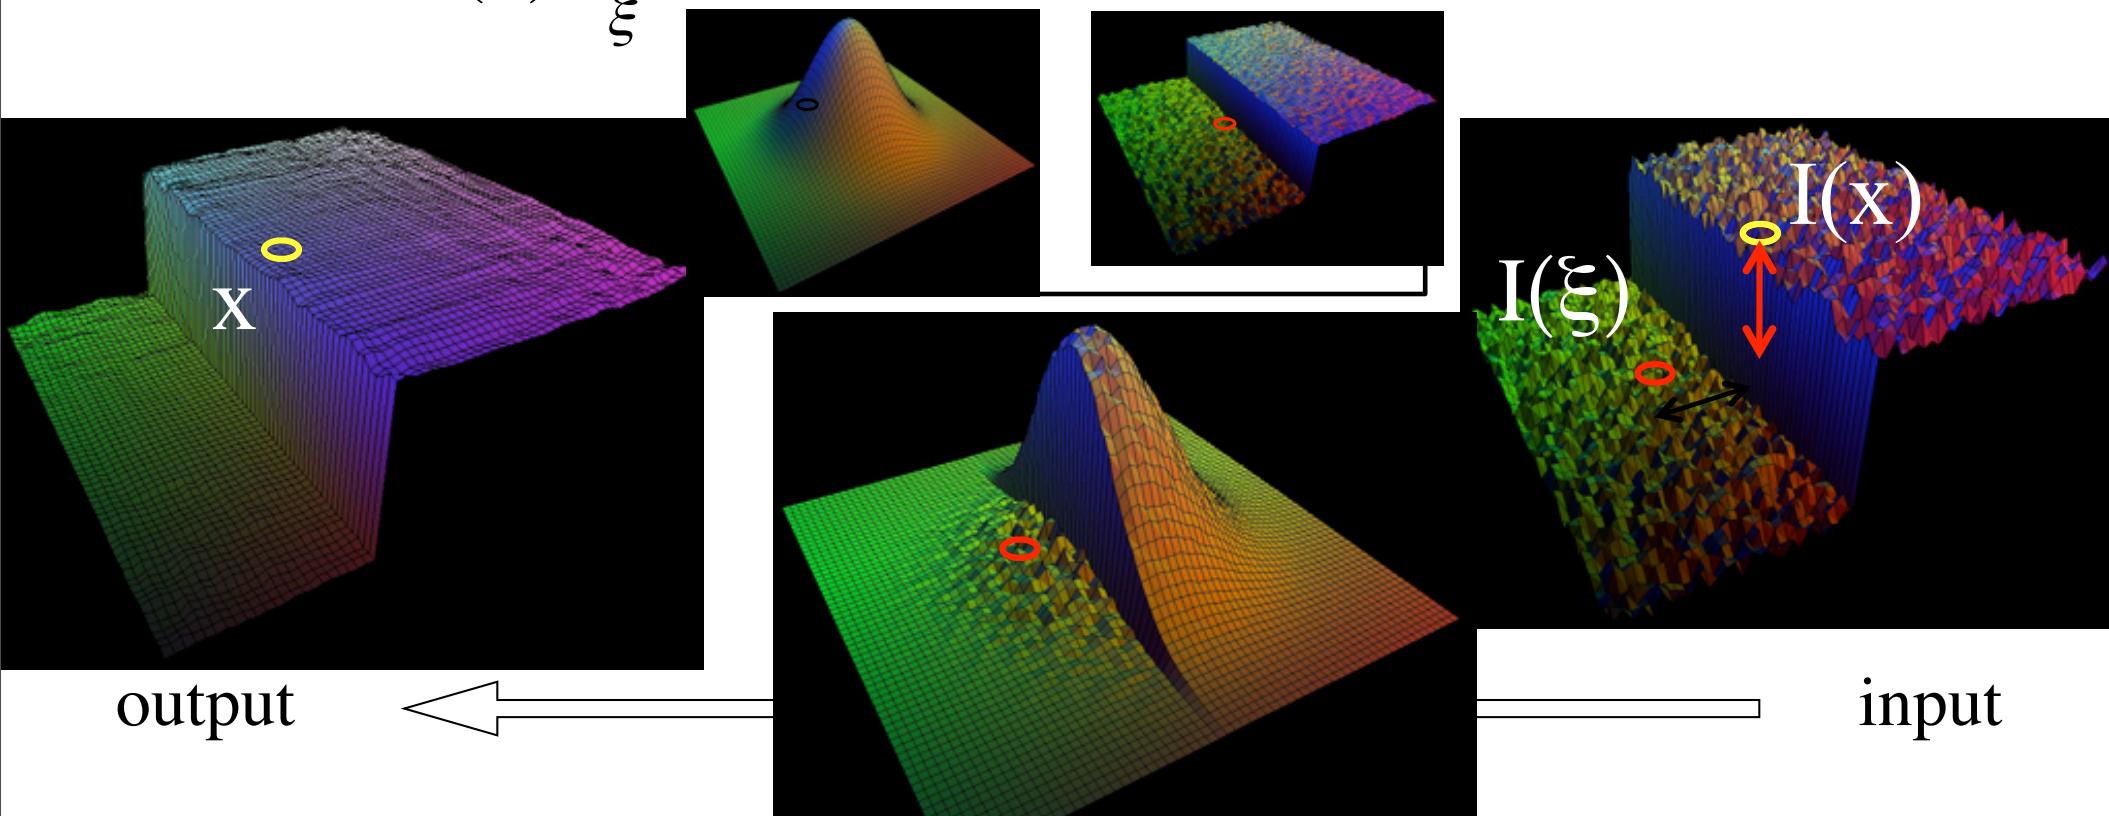
input

Bilateral filtering

[Tomasi and Manduchi 1998]

- Spatial Gaussian f
- Gaussian \mathbf{g} on the intensity difference

$$J(x) = \frac{1}{k(x)} \sum_{\xi} f(x, \xi) g(I(\xi) - I(x)) I(\xi)$$

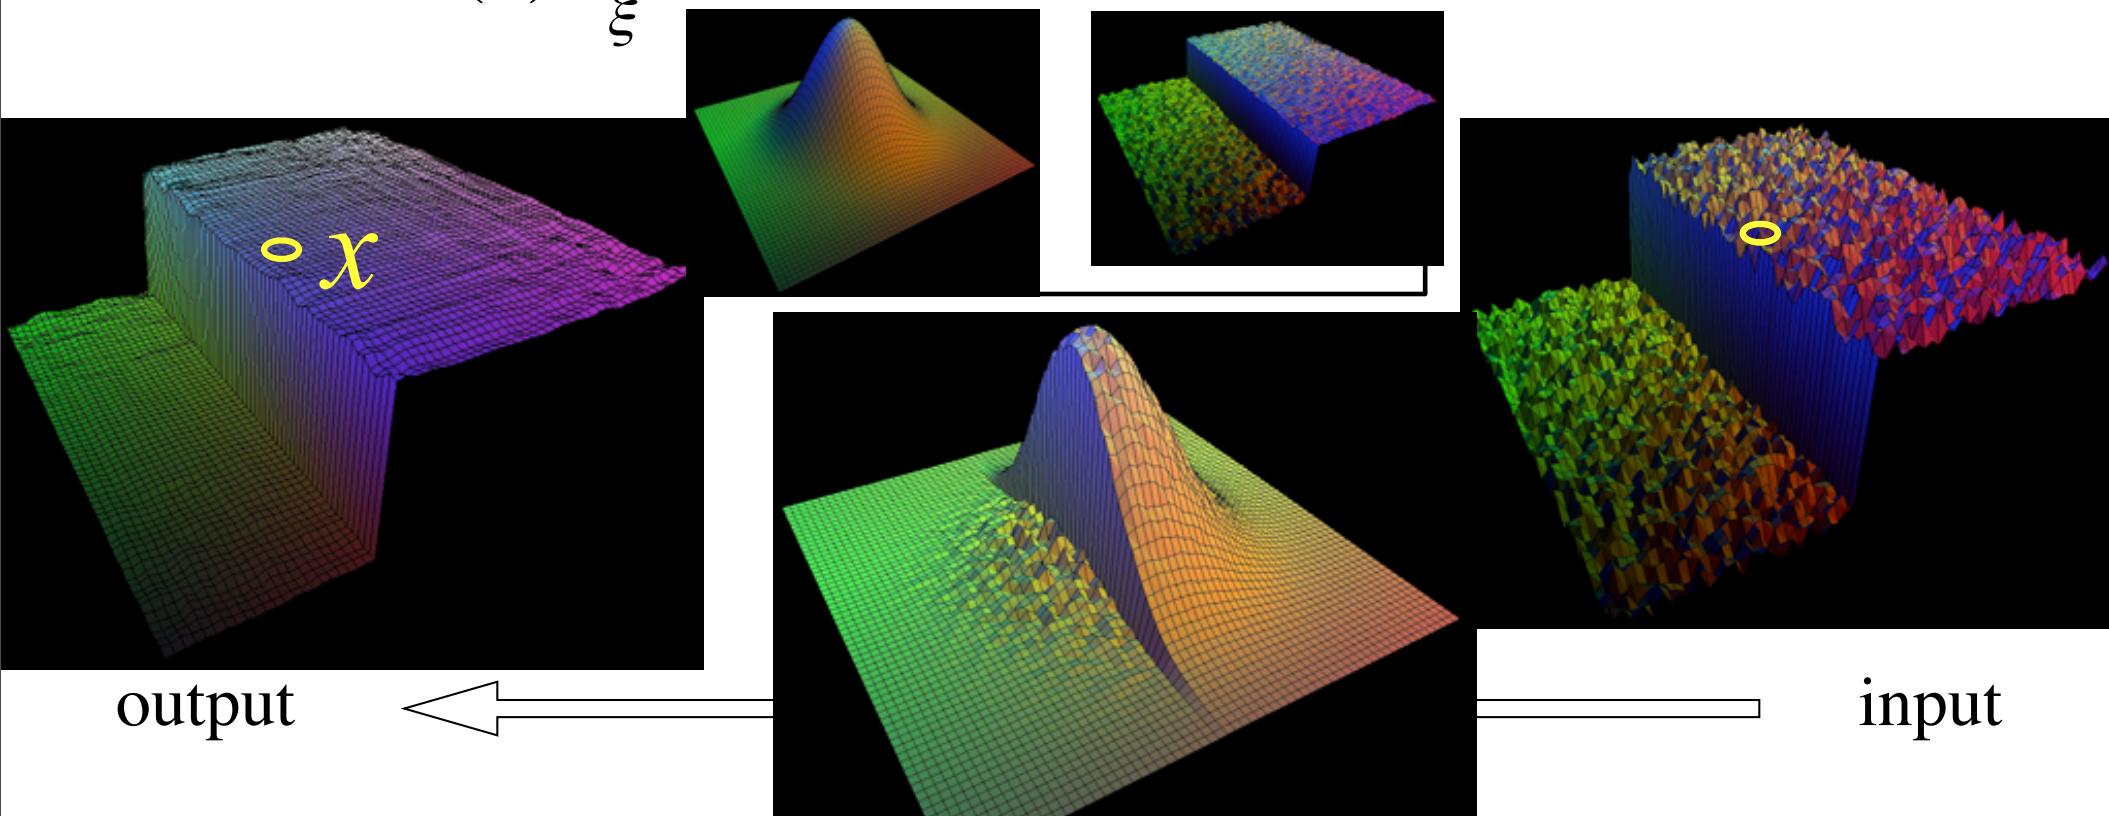


Normalization factor

[Tomasi and Manduchi 1998]

- $k(x) = \sum_{\xi} f(x, \xi) g(I(\xi) - I(x))$

$$J(x) = \frac{1}{k(x)} \sum_{\xi} f(x, \xi) g(I(\xi) - I(x)) I(\xi)$$

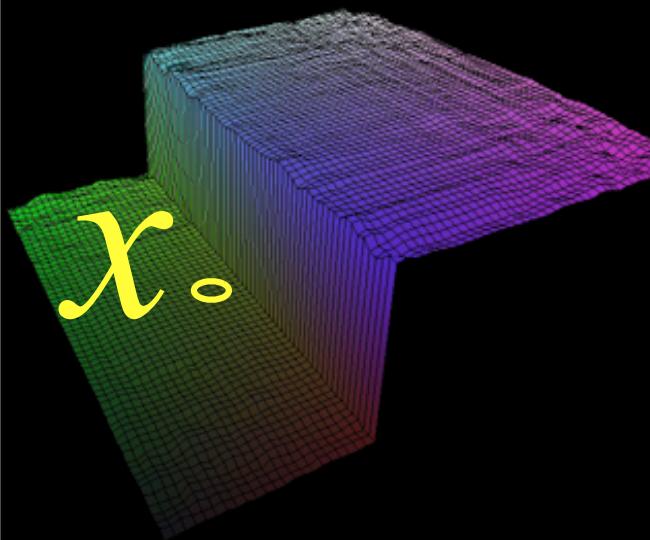


Bilateral filtering is non-linear

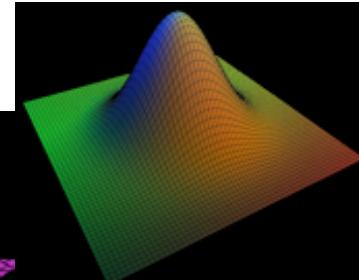
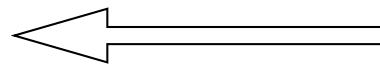
[Tomasi and Manduchi 1998]

- The weights are different for each output pixel

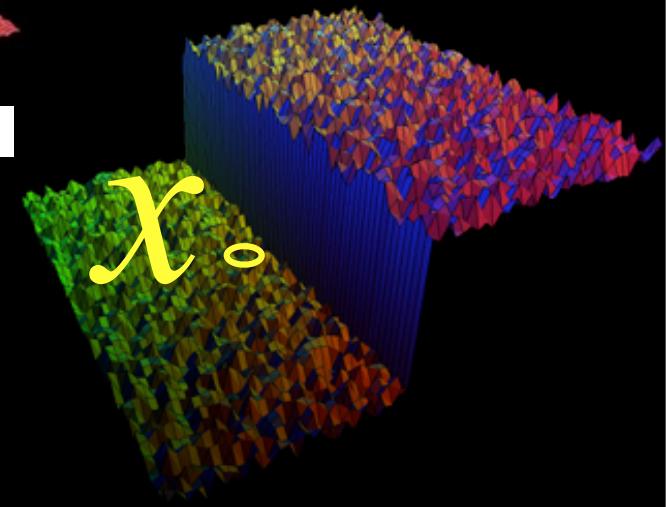
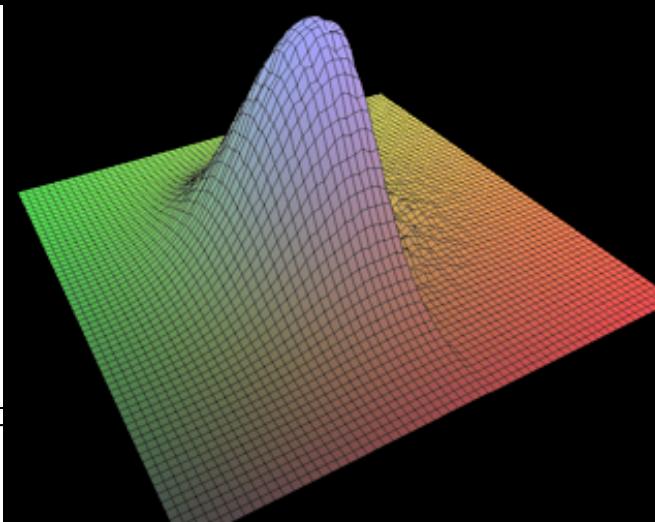
$$J(x) = \frac{1}{k(x)} \sum_{\xi} f(x, \xi) g(I(\xi) - I(x)) I(\xi)$$



output



|

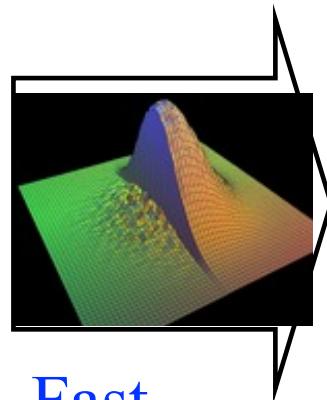


input

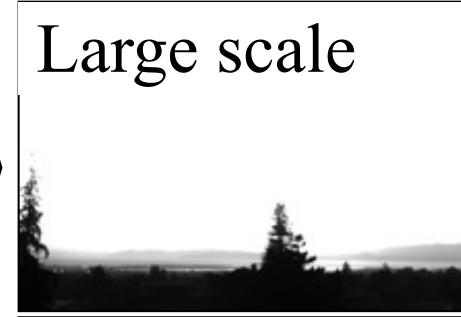
Log domain

- **Very important to work in the log domain**
- **Recall: humans are sensitive to multiplicative contrast**
- **With log domain, our notion of “strong edge” always corresponds to the same contrast**

Recap



Fast
Bilateral
Filter
IN LOG



detail=
input log - large scale

**Reduce
contrast**

Preserve!



