

CSC 433/533 Computer Graphics

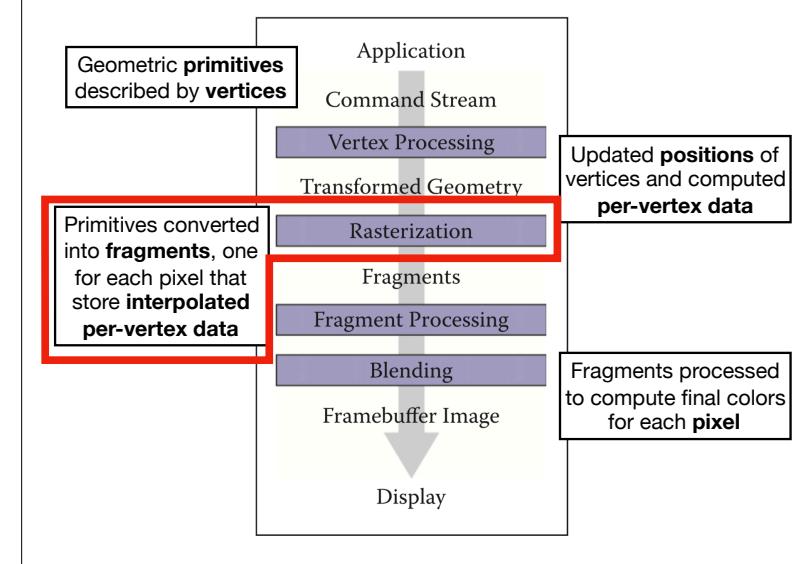
Joshua Levine
josh@email.arizona.edu

Lecture 21 Graphics Pipeline 2

Nov. 13, 2019

Today's Agenda

- Reminders:
 - A05 DUE. A06 posted
- Goals for today: finish discussing the graphics pipeline by talking about the operations before and after rasterization



Recall: Barycentric Coordinates

- A coordinate system to write all points \mathbf{p} as a weighted sum of the vertices

$$\mathbf{p} = \alpha\mathbf{a} + \beta\mathbf{b} + \gamma\mathbf{c}$$

$$\alpha + \beta + \gamma = 1$$

- Equivalently, α, β, γ are the proportions of area of subtriangles relative total area, A

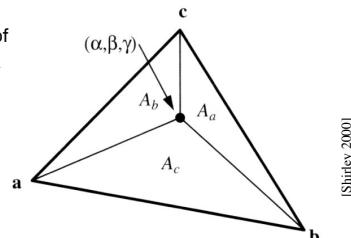
$$A_a / A = \alpha$$

$$A_b / A = \beta$$

$$A_c / A = \gamma$$

- Triangle interior test:

$$\alpha > 0, \beta > 0, \text{ and } \gamma > 0$$



Computing Barycentric Coordinates for Pixels

- We can rely on the equations for each edge of the triangle. Given points $(x_0, y_0), (x_1, y_1)$, and (x_2, y_2) on a triangle, we know:

$$f_{01}(x, y) = (y_0 - y_1)x + (x_1 - x_0)y + x_0y_1 - x_1y_0$$

$$f_{12}(x, y) = (y_1 - y_2)x + (x_2 - x_1)y + x_1y_2 - x_2y_1$$

$$f_{20}(x, y) = (y_2 - y_0)x + (x_0 - x_2)y + x_2y_0 - x_0y_2$$

- And we can write the barycentric coordinate as:

$$\alpha = f_{12}(x, y) / f_{12}(x_0, y_0)$$

$$\beta = f_{20}(x, y) / f_{20}(x_1, y_1)$$

$$\gamma = f_{01}(x, y) / f_{01}(x_2, y_2)$$

Acceleration: #1

- Only check x, y that are in the bounding box of the triangle

```

xmin = floor(min({x_0,x_1,x_2}))
xmax = ceiling(max({x_0,x_1,x_2}))
ymin = floor(min({y_0,y_1,y_2}))
ymax = ceiling(max({y_0,y_1,y_2}))

for all x in [xmin,xmax] {
    for all y in [ymin,ymax] {
        compute ( $\alpha, \beta, \gamma$ ) for (x,y)
        if ( $\alpha \in [0,1]$  and  $\beta \in [0,1]$  and  $\gamma \in [0,1]$ ) {
            draw(x,y);
        }
    }
}
  
```

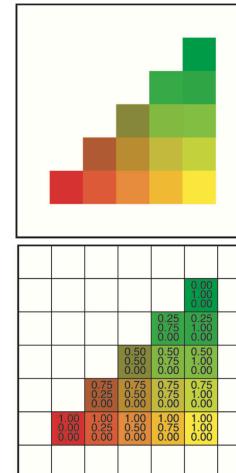
Interpolating Color w/ Gouraud Shading

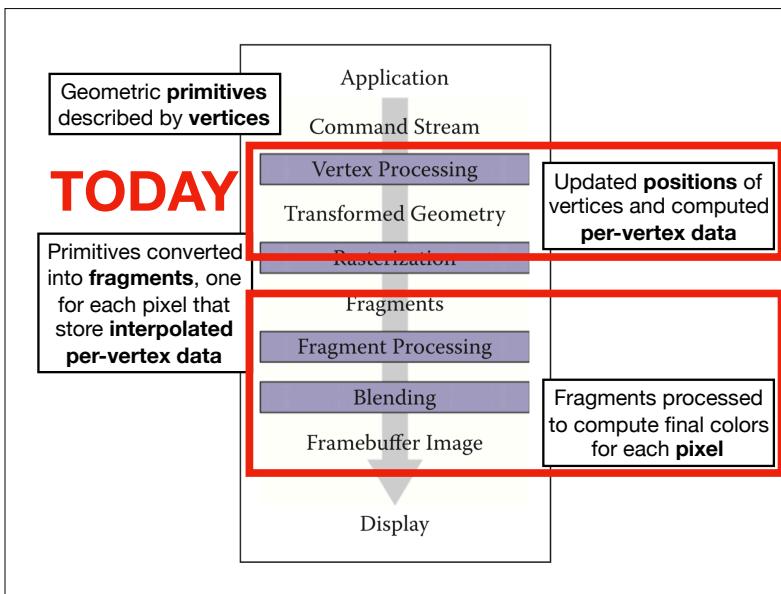
- Having the barycentric coordinates means that we can interpolate color:

$$\mathbf{c} = \alpha\mathbf{c}_0 + \beta\mathbf{c}_1 + \gamma\mathbf{c}_2$$

```

for all x in [xmin,xmax] {
    for all y in [ymin,ymax] {
        compute ( $\alpha, \beta, \gamma$ ) for (x,y)
        if ( $\alpha, \beta, \gamma \in [0,1]$ ) {
            c =  $\alpha\mathbf{c}_0 + \beta\mathbf{c}_1 + \gamma\mathbf{c}_2$ 
            draw(x,y);
        }
    }
}
  
```





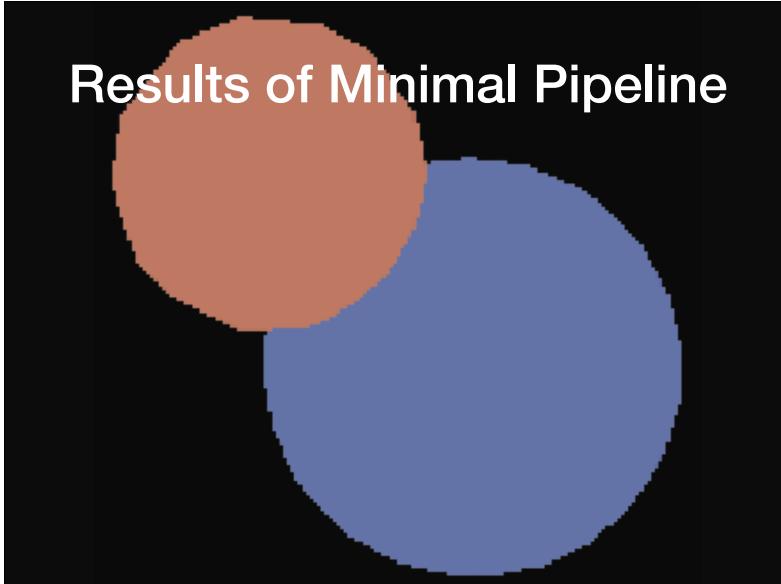
Operations Before and After Rasterization

- Vertex Processing:
 - Role: Prepare the data necessary for rasterization.
- Fragment Processing:
 - Role: Combine the fragments and compute additional information to determine the final color for each pixel.

A Minimal Pipeline for Raster Graphics

- Vertex Processing
 - Transform vertex positions from object to screen space
 - Set color for each vertex based on primitive.
- Rasterizer
 - Enumerate fragments for each primitive and set their color.
- Fragment Processing
 - Write fragment colors to the appropriate pixel

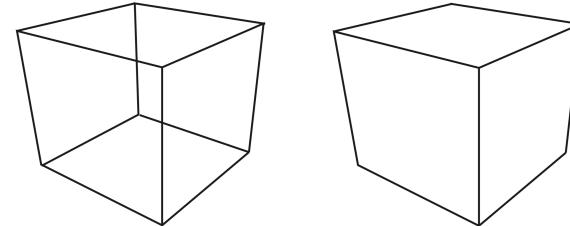
Results of Minimal Pipeline



Resolving Issues with Depth

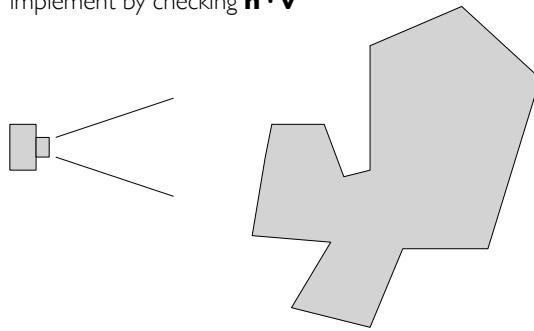
Hidden Surface Removal

- Perspective projection provides one important cue for 3D relationships
- Depth / Occlusion is equally important



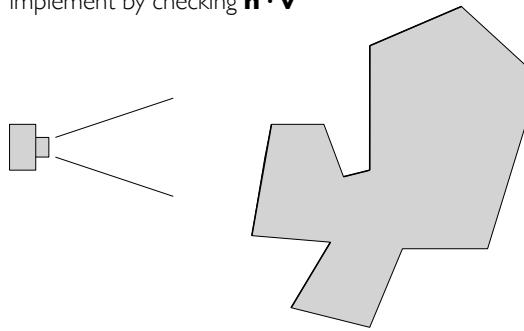
Back face culling

- **For closed shapes you will never see the inside**
 - therefore only draw surfaces that face the camera
 - implement by checking $\mathbf{n} \cdot \mathbf{v}$



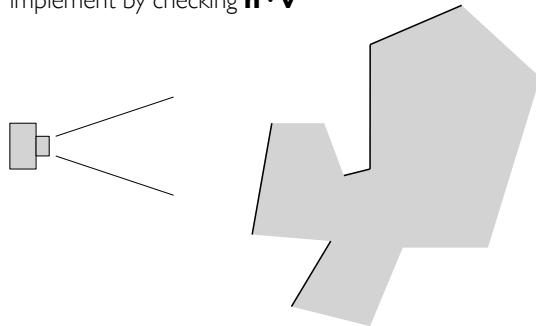
Back face culling

- **For closed shapes you will never see the inside**
 - therefore only draw surfaces that face the camera
 - implement by checking $\mathbf{n} \cdot \mathbf{v}$



Back face culling

- **For closed shapes you will never see the inside**
 - therefore only draw surfaces that face the camera
 - implement by checking $\mathbf{n} \cdot \mathbf{v}$

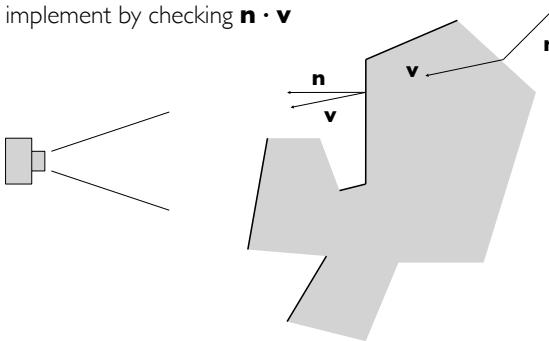


Cornell CS4620 Spring 2018 • Lecture 11

© 2018 Steve Marschner • 5

Back face culling

- **For closed shapes you will never see the inside**
 - therefore only draw surfaces that face the camera
 - implement by checking $\mathbf{n} \cdot \mathbf{v}$

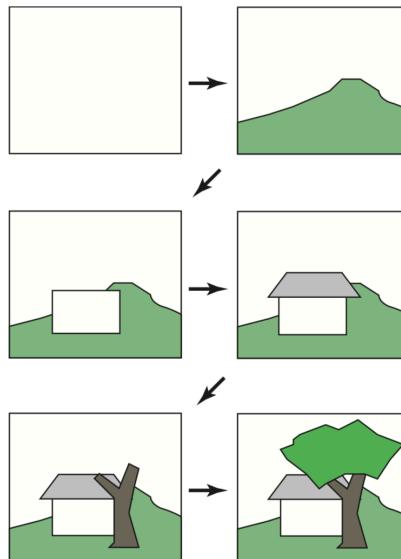


Cornell CS4620 Spring 2018 • Lecture 11

© 2018 Steve Marschner • 5

Painter's Algorithm

- Simple way to do hidden surfaces
- Draw from back-to-front, overwriting directly on the image



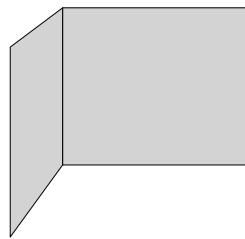
Painter's Algorithm

- Draw one primitive at a time.



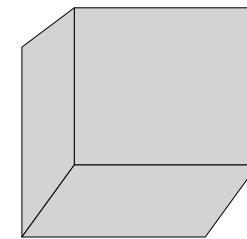
Painter's Algorithm

- Draw one primitive at a time.



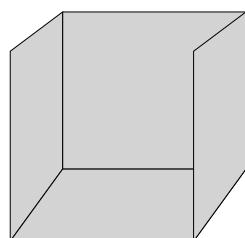
Painter's Algorithm

- Draw one primitive at a time.



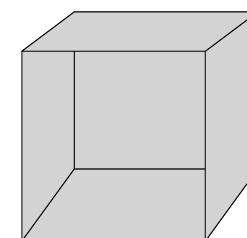
Painter's Algorithm

- Draw one primitive at a time.



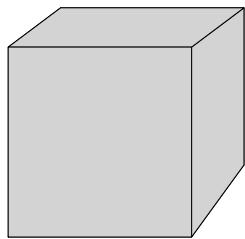
Painter's Algorithm

- Draw one primitive at a time.



Painter's Algorithm

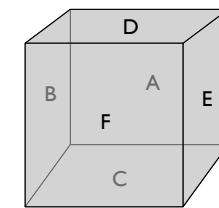
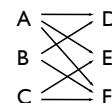
- Draw one primitive at a time.



Painter's algorithm

- **Amounts to a topological sort of the graph of occlusions**

- that is, an edge from A to B means A sometimes occludes B
- any sort is valid
 - ABCDEF
 - BADCFE
- if there are cycles
there is no sort



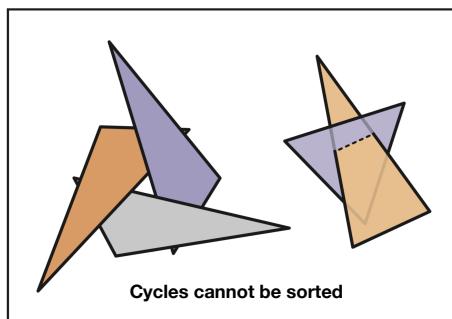
Cornell CS4620 Spring 2018 • Lecture 11

© 2018 Steve Marschner • 7

Painter's algorithm

- **Amounts to a topological sort of the graph of occlusions**

- that is, an edge from A to B means A sometimes occludes B
- any sort is valid
 - ABCDEF
 - BADCFE
- if there are cycles
there is no sort



Using a z-Buffer for Hidden Surfaces

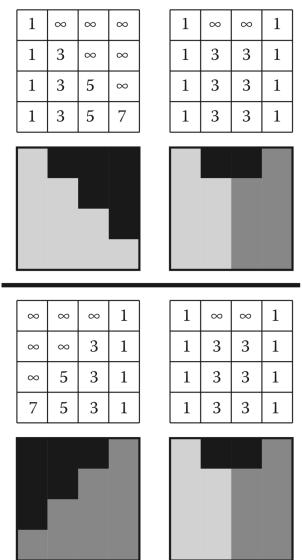
- Most of the time, sorting the primitives in z is too expensive and complex
- Solution: draw primitives in any order, but keep track of closest at the fragment level
 - Method: use an extra data structure that tracks the closest fragment in depth.
- When drawing, compare fragment's depth to closest current depth and discard the one that is further away

Cornell CS4620 Spring 2018 • Lecture 11

© 2018 Steve Marschner • 7

z-Buffers

- Example w/ drawing two triangles
- Order doesn't matter
- Often depths are stored as integers to keep memory overhead low
- Memory-intensive, brute force approach, but it works and it is the standard because of its simplicity



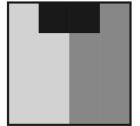
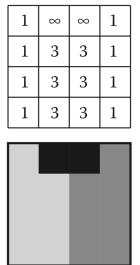
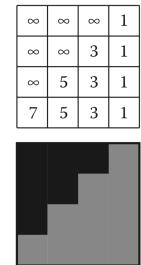
The diagram illustrates the concept of z-buffers. It shows four tables representing depth values for a 2x4 pixel area, followed by two grayscale images showing the resulting rendered output.

1	∞	∞	∞
1	3	∞	∞
1	3	5	∞
1	3	5	7

1	∞	∞	1
1	3	3	1
1	3	3	1
1	3	3	1

∞	∞	∞	1
∞	∞	3	1
∞	5	3	1
7	5	3	1

1	∞	∞	1
1	3	3	1
1	3	3	1
1	3	3	1

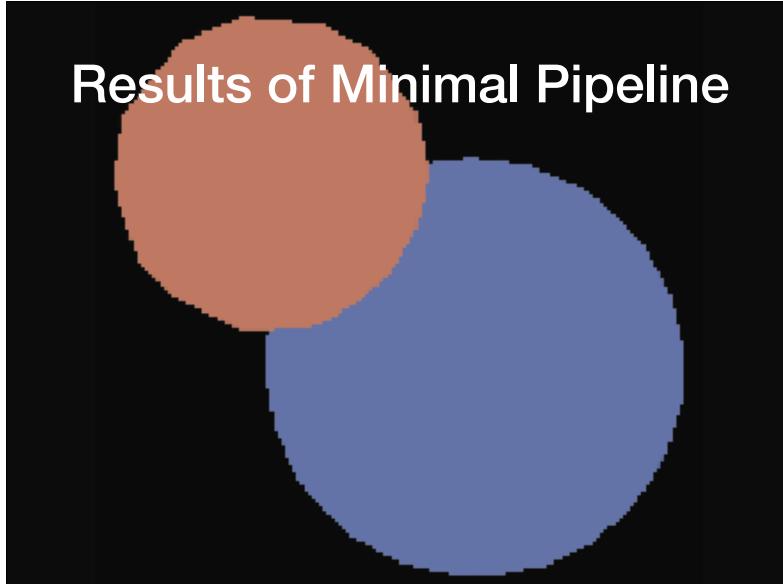
z-Buffer Precision

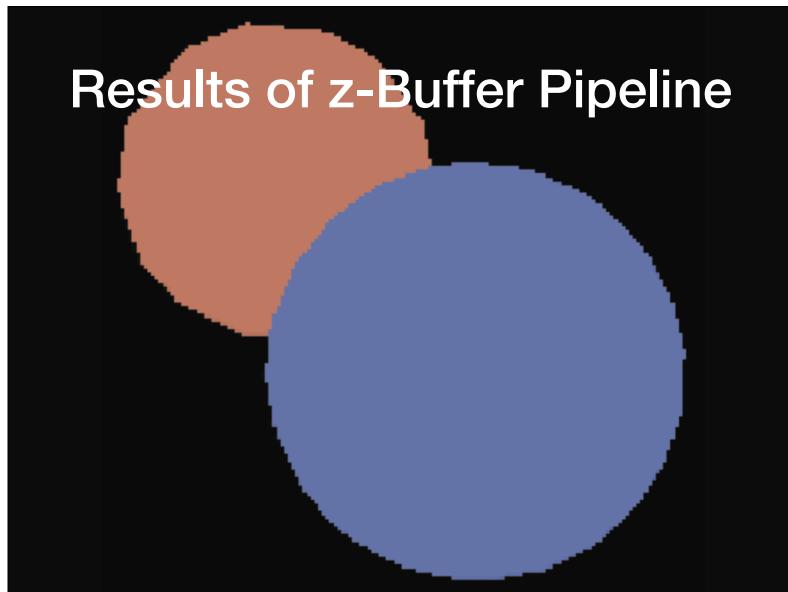
- If using integers, the precision is related to the distance between the near and far clipping planes
 - Coincidentally, this necessitates the existence of the near/far planes (unlike with ray tracing)
- Generally, use z values that are after transformation
- More info in the book on this topic!

Pipeline for z-Buffer

- Vertex Processing
 - Transform vertex positions, set color for each vertex based on primitive.
- Rasterizer
 - Enumerate fragments for each primitive and set their color.
 - Interpolate the z-value for each fragment
- Fragment Processing
 - Write fragment colors only if interpolated z is closer

Results of Minimal Pipeline





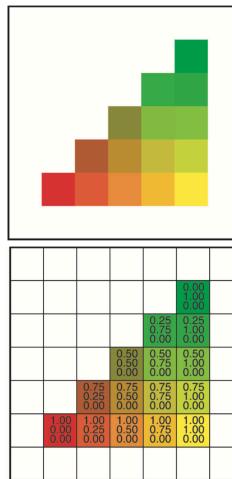
Resolving Issues with Interpolation

Recall: Screen Space Interpolation of Color

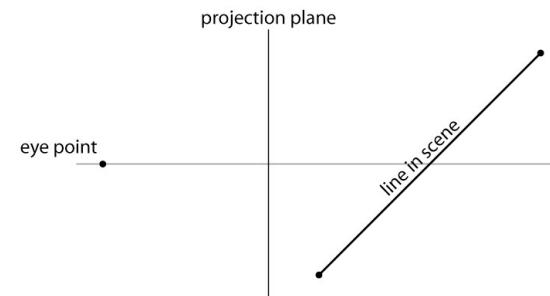
- Having the barycentric coordinates means that we can interpolate color:

$$\mathbf{c} = \alpha\mathbf{c}_0 + \beta\mathbf{c}_1 + \gamma\mathbf{c}_2$$

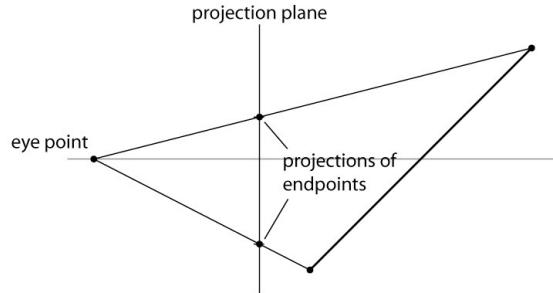
```
for all x in [xmin,xmax] {
    for all y in [ymin,ymax] {
        compute ( $\alpha$ ,  $\beta$ ,  $\gamma$ ) for (x,y)
        if ( $\alpha, \beta, \gamma \in [0,1]$ ) {
            c =  $\alpha\mathbf{c}_0 + \beta\mathbf{c}_1 + \gamma\mathbf{c}_2$ 
            draw(x,y);
        }
    }
}
```



Interpolating in projection

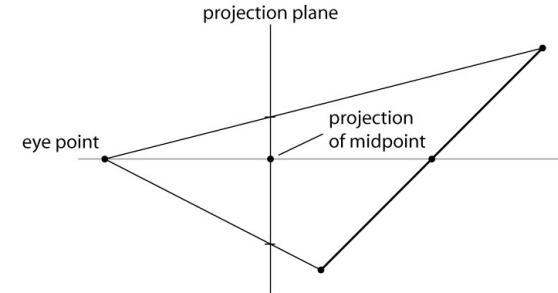


Interpolating in projection



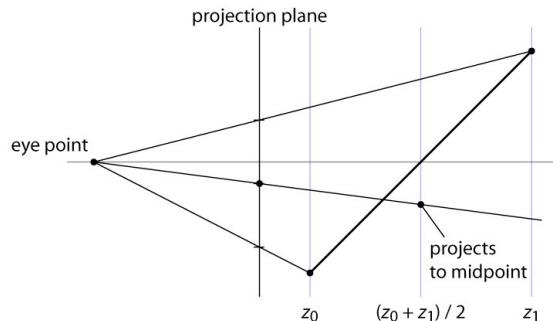
linear interp. in screen space \neq linear interp. in world (eye) space

Interpolating in projection



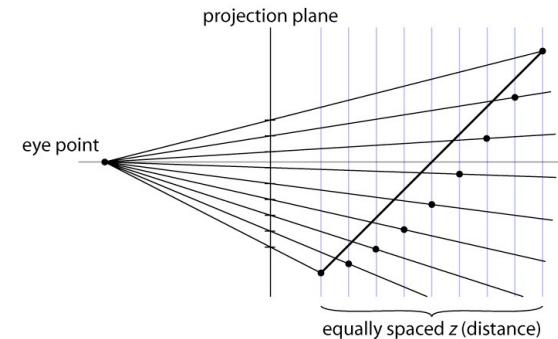
linear interp. in screen space \neq linear interp. in world (eye) space

Interpolating in projection



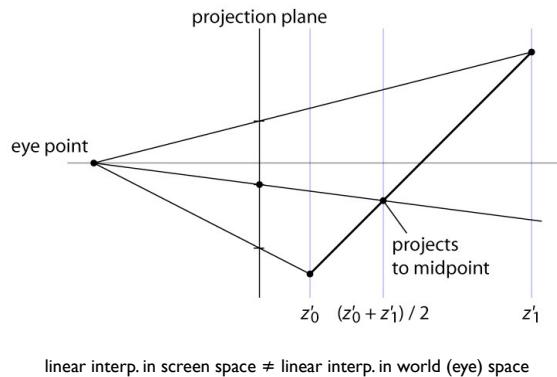
linear interp. in screen space \neq linear interp. in world (eye) space

Interpolating in projection



linear interp. in screen space \neq linear interp. in world (eye) space

Interpolating in projection

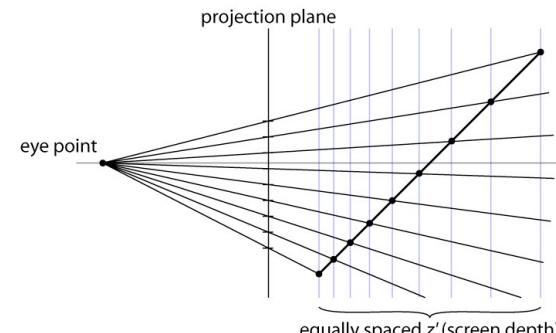


linear interp. in screen space \neq linear interp. in world (eye) space

Cornell CS4620 Spring 2018 • Lecture 11

© 2018 Steve Marschner • 12

Interpolating in projection



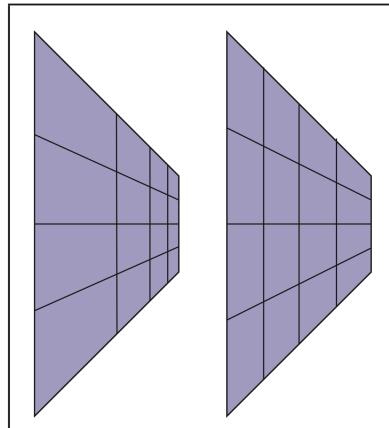
linear interp. in screen space \neq linear interp. in world (eye) space

Cornell CS4620 Spring 2018 • Lecture 11

© 2018 Steve Marschner • 12

Screen Space vs. World Space

- Linear variation in screen space does not yield linear variation in world space due to projection
- Consider a checkerboard at a steep angle; all squares are the same size on the plane, but not on screen



Barycentric Interpolation in World Space

- We need to account for the homogenous rescaling in any function we interpolate:

```
for all x in [xmin,xmax] {
    for all y in [ymin,ymax] {
        compute ( $\alpha$ ,  $\beta$ ,  $\gamma$ ) for (x,y)
        if ( $\alpha, \beta, \gamma \in [0,1]$ ) {
            z =  $\alpha z_0/w_0 + \beta z_1/w_1 + \gamma z_2/w_2$ 
            s =  $\alpha/w_0 + \beta/w_1 + \gamma/w_2$ 
            fragment.z = z/s
        }
    }
}
```

- Section 11.2.4 of the textbook describes this in more detail

Adding 3D Cues w/ Shading and Lighting

Flat Shading

- Use the normal of the primitive (triangle) to compute a shaded color based on the lights in the scene
- Creates a faceted appearance that highlights the mesh geometry

Pipeline for Flat Shading

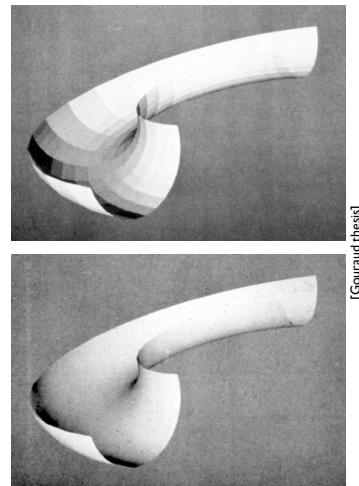
- Vertex Processing
 - Compute shaded color per triangle using triangle normal
 - Transform vertex positions
- Rasterizer
 - Interpolate the z-value for each fragment
 - Set fragment color based on its primitive color.
- Fragment Processing
 - Write fragment colors only if interpolated z is closer

Results of Flat Shading Pipeline



Gouraud shading

- Often we're trying to draw smooth surfaces, so facets are an artifact
 - compute colors at vertices using vertex normals
 - interpolate colors across triangles
 - "Gouraud shading"
 - "Smooth shading"



Cornell CS4620 Spring 2017 • Lecture 13

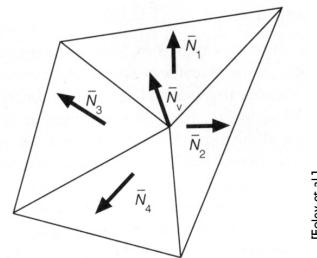
© 2017 Steve Marschner • 21

Pipeline for Gouraud Shading

- Vertex Processing
 - Compute shaded color per vertex using vertex normal
 - Transform vertex positions
- Rasterizer
 - Interpolate the z-value for each fragment as well as r,g,b colors
- Fragment Processing
 - Write fragment colors only if interpolated z is closer

Vertex normals

- Need normals at vertices to compute Gouraud shading
- Best to get vtx. normals from the underlying geometry
 - e. g. spheres example
- Otherwise have to infer vtx. normals from triangles
 - simple scheme: average surrounding face normals



[Foley et al.]

$$N_v = \frac{\sum_i N_i}{\|\sum_i N_i\|}$$

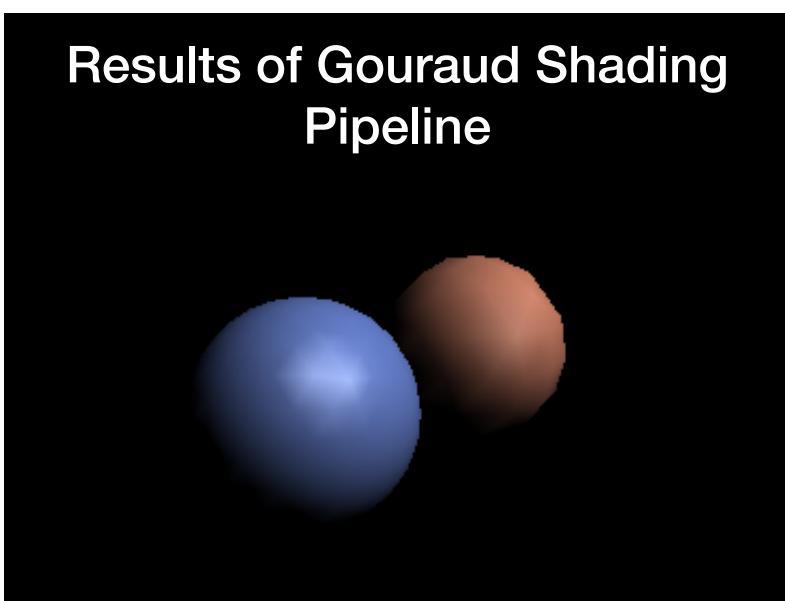
Cornell CS4620 Spring 2018 • Lecture 11

© 2018 Steve Marschner • 28

Results of Flat Shading Pipeline



Results of Gouraud Shading Pipeline



Problems w/ Gouraud Shading

- While you can use any shading model on the vertices (Gouraud shading is just an interpolation method!), typically using only diffuse color works best.
- Results tend to be poor with rapidly-varying models like specular color
 - In particular, when triangles are large relative to how quickly color is changing

Per-pixel (Phong) shading

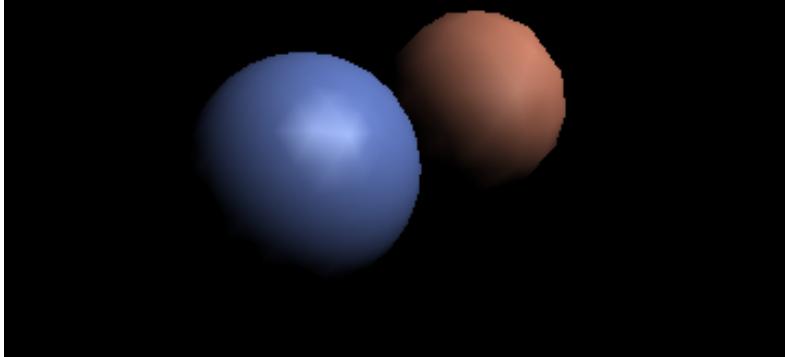
- **Get higher quality by interpolating the normal**
 - just as easy as interpolating the color
 - but now we are evaluating the illumination model per pixel rather than per vertex (and normalizing the normal first)
 - in pipeline, this means we are moving illumination from the vertex processing stage to the fragment processing stage



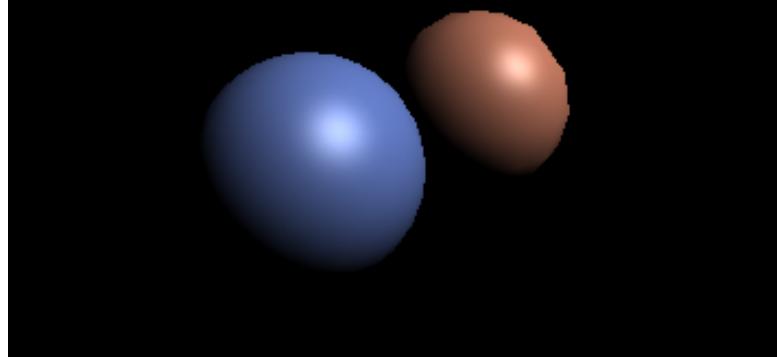
Pipeline for Phong Shading

- Vertex Processing
 - Transform vertex positions, compute normal for vertex based on primitive.
- Rasterizer
 - Interpolate world space position and normal of each fragment, and store the barycentric weights
- Fragment Processing
 - Compute shading using position and normal and fixed color for each primitive.
 - Write fragment colors only if interpolated z is closer

Results of Gouraud Shading Pipeline



Results of Phong Shading Pipeline

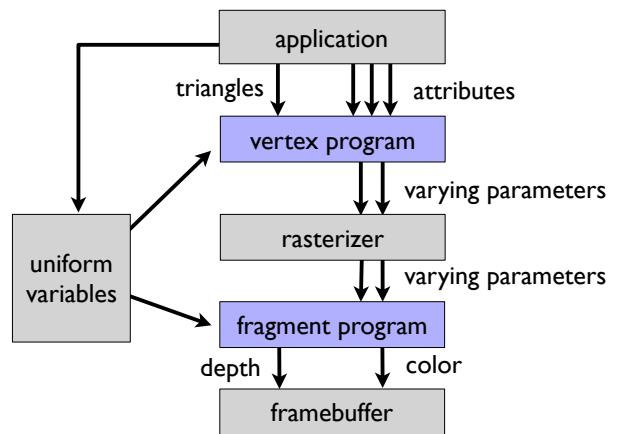


Some Comments on Hardware

Programming hardware pipelines

- **Modern hardware graphics pipelines are flexible**
 - programmer defines exactly what happens at each stage
 - do this by writing *shader programs* in domain-specific languages called *shading languages*
 - rasterization is fixed-function, as are some other operations (depth test, many data conversions, ...)
- **One example: OpenGL and GLSL (GL Shading Language)**
 - several types of shaders process primitives and vertices; most basic is the *vertex program*
 - after rasterization, fragments are processed by a *fragment program*

GLSL Shaders



Cornell CS4620 Spring 2018 • Lecture 11

© 2018 Steve Marschner • 35

Lec22 Required Reading

- FOG, Ch. 11.1-11.14

Reminder: Assignment 06

Assigned: Wednesday, Nov. 13
Written Due: Monday, Nov. 25, 4:59:59 pm
Program Due: Wednesday, Nov. 27, 4:59:59 pm